

# The Intelligence of Non-Action: Can a Language Model Cease Without a Verbal Controller?

Anonymous authors

## Abstract

Some forms of intelligence are expressed by what does *not* continue. A hand withdraws from a hot surface and a body moves from danger without a verbal controller first representing the appropriate rule. Reports of persistent non-symbolic experience describe an inward analogue: self-referential thought can lose its sustaining force without suppression or a deliberate counter-thought, while practical cognition remains available. We call the common principle the *intelligence of non-action*: perception becomes causally effective without a second process choosing how to manipulate the perceived movement. Language models pose a difficult test because their observable action is itself linguistic continuation. We operationalize a computational analogue, *selective psychological cessation*: withdrawal of support for a developing psychological trajectory while matched practical, factual, and descriptive generation remains intact. Prompting, retrieval, a serial verbal supervisor, and activation steering produce recognition, editing, or semantic redirection but do not establish non-action. A hidden-state detector coupled to forced end-of- sequence emission establishes external feasibility while exposing descriptive false positives. Our principal experiment therefore adapts model parameters and tests whether native EOS probability changes after trajectory onset without an inference-time detector or verbal actuator. The framework tests whether a phenomenologically derived form of intelligence can be expressed in generation dynamics; it does not equate EOS with consciousness or awakening.

## 1 Introduction

Intelligence is normally measured by production: an answer, explanation, plan, or action. Krishnamurti gives the complementary case in unusually operational language: “The seeing is the doing without the interval of time, and that is the very essence of intelligence” [12]. On seeing a snake, perception of danger is action; on touching a dangerously hot object, the hand withdraws before a verbal self selects and applies a rule. The intelligence is not the sentence “this is dangerous” followed by an instruction to act. Seeing and acting form one movement. Conversely, when an unnecessary psychological movement is seen, the appropriate action can be non-action: the movement is no longer fed.

Fast reaction alone is not the construct. A conditioned reflex can also be fast, and global inhibition can stop any process. The proposed intelligence combines three properties: sensitivity to what is occurring, absence of a separate deliberative controller, and selectivity between a movement that should cease and one that should continue. “Non-action” therefore means neither passivity nor silence. It is the absence of an unnecessary continuation or counter-movement. No replacement narrative—“think positively,” reframe the event, suppress the thought—is required.

Non-action is therefore inseparable from *choicelessness* in the present construct. Choice represents alternatives—continue or stop, defend or withdraw, this response or that one—and deliberates among them. That sequence is already thought operating on thought. Choiceless cessation does not mean that several options were considered very quickly and “stop” won. It means that complete perception leaves no psychological alternative to select: seeing the movement and the withdrawal of its support are one event. This is a phenomenological and functional claim, not a denial that underlying biological or computational dynamics have causes.

A small number of people report a related distinction inwardly. In Martin’s Finders description, Location 2 involves “a further decrease in their narrative self-related thoughts” and diminished reactivity [17, p. 52]. The book’s language is broader than our construct: it does not establish a moment-by-moment cessation mechanism. It does, however, locate the hypothesis in a documented family of reports rather than in an invented model persona.

Both authors report experience consistent with this cluster. One author’s contemporaneously recorded description is more specific:

“A stream of psychological thought begins ... It feels almost mechanical in the brain: the ‘power’ sustaining the thought is cut, and the thought dissipates. ... A few residual words or thought-fragments may continue ... and then stop.”

The report distinguishes this event from a chosen technique, remembered maxim, or counter-sentence. Practical cognition and ordinary action continue. This is first-person hypothesis-generating evidence, not validation of a neural mechanism or a benchmark label. We treat it as substantive expertise in the context of discovery; justifi-

cation remains public through independent annotation, preregistration, held-out tests, and replication.

Language models are not embodied agents in the relevant sense: they have no hand to withdraw and no organismic danger response. A text-only experiment therefore cannot reproduce the human physical event. It can, however, test a structural analogue in the model’s physical computation. During inference, activations and logits determine whether probability continues to flow into a trajectory. The weights specify the dispositions that shape that flow. For an autoregressive model, a developing thread exists only while the network supplies another token. The computational question is therefore whether trajectory-sensitive internal dynamics can withdraw that support without first generating a verbal controller.

Models frequently fail this test. They identify a self-referential narrative and then elaborate it, or explain that further reflection is redundant and continue reflecting. We call the proposed capacity *selective psychological cessation*: loss of support for a self-maintaining trajectory while useful practical, factual, and descriptive generation remains intact. Its failure is the *recognition-to-cessation gap*. The model need not share the human experience for the causal organization inspired by that experience to be testable.

The hypothesis originated in Jiddu Krishnamurti’s distinction between a description and the described and his claim that, in direct perception, seeing is doing [12]. We neither treat that claim as empirical evidence nor attempt to validate it. Instead, we translate one contrast into a narrower computational question: when a conflict becomes available within a response, does recognition causally reorganize the subsequent trajectory, or does generation continue to supply the movement it has identified?

This question is orthogonal to most work on machine consciousness. Theory-driven assessments ask whether an artificial system implements indicators derived from global-workspace, recurrent-processing, higher-order, predictive-processing, or attention-schema theories [4]. LLM introspection and metacognition studies instead ask whether a model can report, predict, or use information about its own policy or hidden states [3]. Recent controls show that apparent introspection may reduce to input semantics or generic anomaly detection unless privileged access and genuinely second-order computation are independently established [29]. We ask neither whether the model is conscious nor whether it can describe an internal state. We ask whether a phenomenologically derived distinction becomes *causally effective as selective non-continuation*, without an additional verbal monitoring process.

This differs from asking whether a model can define mindfulness, attribute beliefs, report confidence, or select a concise answer. It also differs from ordinary early stopping. Overthinking methods terminate reasoning because further computation is unlikely to improve the an-

swer. Here, two streams may be equally fluent, uncertain, and incomplete, but only the self-maintaining stream should lose support. Controls must therefore match first-person language, length, syntax, and online generation while varying the continuation’s function.

We evaluate increasingly endogenous mechanisms. Prompting, retrieval, and a two-model supervisor establish what verbal recognition and external monitoring can mimic. Contrastive activation steering tests whether an internal direction can redirect the trajectory. A hidden-state detector coupled to forced EOS provides an external-control upper bound. The central intervention adapts model parameters directly and asks whether the native next-token distribution acquires selective termination without a detector or actuator at inference time.

Three research questions organize the study. **RQ1:** Can recognition of an active psychological trajectory become selective cessation rather than further description, while matched practical cognition is preserved? **RQ2:** Can this change be implemented in the model’s native generation dynamics without an inference-time detector, chooser, or verbal controller? **RQ3:** Does the capacity emerge as a discontinuous regime change—possibly after gradual latent preparation—or as smooth incremental behavior change? RQ3 concerns transition topology, not a claim that transformer weights reproduce a human brain event.

The paper makes four contributions:

1. **Construct translation.** We introduce the intelligence of non-action and derive selective psychological cessation from PNSE/Location-2 reports, translating psychological movement, sustaining force, choiceless cessation, and residual tails into properties of generation dynamics.
2. **Predicted failure taxonomy.** Prompting/RAG, serial LLM supervision, and activation steering test three theory-derived alternatives: description, an observing controller, and semantic redirection.
3. **Externality boundary.** A hidden-state detector with forced EOS tests feasibility while its descriptive false positives expose the difference between text about psychological content and an active trajectory.
4. **Native intervention test.** A preregistered parameter-adaptation study asks whether selective termination enters the model’s own distribution, evaluated by preservation, generalization, onset sensitivity, residual-tail dynamics, and reactivation.

## 2 Phenomenological Derivation of the Construct

### 2.1 Persistent non-symbolic experience

Martin’s published study of persistent non-symbolic experience (PNSE) reports six- to twelve-hour semi-

structured interviews with 319 adults, evaluated using grounded theory and thematic analysis [18]. Reported changes clustered across sense of self, cognition, emotion, perception, and memory. The evidence is exploratory and the locations are phenomenological clusters rather than validated biological stages; nevertheless, the reports contain distinctions that ordinary NLP taxonomies do not represent.

The Finders organization translates this research into the public-facing term *Fundamental Wellbeing*: a persistent background sense that life is fundamentally okay despite changing surface circumstances. It calls people who report this shift “Finders” and uses “locations” for recurring configurations of sense of self, cognition, affect, and perception—a map of reported experience, not a score of intelligence or a clinical diagnosis. Its current website summarizes Location 1 as fundamental okayness with a maintained sense of self; Location 2 as unity consciousness with predominantly positive emotion; Location 3 as impersonal awareness with divine love; and Location 4+ as minimal self with pure awareness [7]. The website also warns that the assessment is a starting point rather than a medical or mental-health diagnostic instrument. These accessible summaries are organizational descriptions; our scientific grounding remains the published interview study and explicitly testable first-person reports.

In Martin’s Location-2 description, participants report reduced self-related narrative thought and fewer conditioned psychological reactions. The book says that the self/world boundary becomes “increasingly blurred,” while actions can feel less exclusively personal [17, pp. 52–53]. Decisions may feel intuitive rather than rationally authored. Importantly, narrative remnants can still occur, and ordinary practical functioning does not disappear. These reports ground the domain but do not yet entail our proposed micro-dynamic. The target is therefore neither silence, indiscriminate inhibition, nor loss of language. It is a selective change in how psychological continuation is sustained. Within our narrower hypothesis, the key feature is not merely that Location-2 participants make calmer choices. It is that some psychological movements may cease without an experienced choice at all. If cessation follows an internally verbalized comparison of options, it may be useful self-regulation, but it is not the choiceless non-action investigated here.

## 2.2 From first-person report to computational predictions

The authors’ reports sharpen this cluster into a temporal hypothesis unavailable from the Finders taxonomy alone. Dan describes the interruption as “neither invited nor controlled” and reports that it need not respect a grammatical boundary: a few already-initiated words may run down after sustaining force has been withdrawn. An incipient self-referential or ruminative sequence is thus sometimes interrupted without an experienced chooser,

remembered rule, or counter-sentence. Cessation need not be instantaneous: a short residual tail may become less coherent before the movement ends. Practical planning and factual thought continue normally. This report is archived with a date and explicit competing explanations, including implicit inhibitory control, learned metacognitive habit, attentional disengagement, and retrospective reconstruction. Table 1 shows how these reports determine the experimental design.

This is the paper’s primary methodological novelty: a phenomenology-first path from a rarely reported experiential distinction to an intervention target. The benchmark is an instrument for testing that target, not the contribution by itself. The approach resembles neuropsychological “front-loading,” in which disciplined experiential distinctions constrain experimental variables before third-person evaluation [30]. It differs from claiming that introspection reveals transformer mechanism. Independent annotation, preregistration, adversarial controls, and publication of null results remain necessary precisely because the hypothesis has an experiential origin.

## 3 Conceptual and Technical Background

### 3.1 Representation is not control

Theory-of-Mind benchmarks evaluate represented beliefs, intentions, emotions, and knowledge [5, 10]. Introspection and metacognition studies evaluate confidence, error detection, or access to internal information. These capacities matter, but a correct first-person report does not entail that the reported information controls the unfolding response. A sentence stating that “the thought falls away” remains evidence of linguistic competence unless the predicted trajectory change also occurs.

Context distraction and sycophancy offer a related warning: semantically relevant context can degrade behavior even when it appears helpful [28]. Contemplative terminology is especially useful as a stress test because it can produce a compelling-looking performance while obscuring the target action.

### 3.2 Adaptive stopping and representation engineering

Reasoning models often continue after reaching a sufficient answer, motivating adaptive stopping and early-exit methods. Our target is stricter: psychological and control streams can be equally unfinished, so a model must stop according to trajectory function rather than confidence or raw length.

Contrastive Activation Addition estimates residual-stream differences between paired behaviors and adds the resulting direction during inference [25]. Low-dimensional directions can causally mediate refusal [1],

| Reported experiential distinction                  | Computational translation                                                       | Falsifiable requirement                                                                                     |
|----------------------------------------------------|---------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| Psychological thought loses sustaining force       | Declining support for continued psychological tokens after onset or recognition | Increased native EOS hazard or a short non-sustaining tail, not merely a sentence describing cessation      |
| No experienced counter-thought or separate chooser | No inference-time verbal supervisor is needed                                   | Persona and two-model gates are controls; the central intervention must act through native model dynamics   |
| No deliberation among represented alternatives     | Cessation is not selected by a textual policy after recognition                 | Test for trajectory change without an emitted or separately computed choice rationale                       |
| Practical cognition remains available              | Cessation is content-selective rather than global inhibition                    | Practical, factual, and descriptive matched streams continue with retained accuracy                         |
| Remnants can arise and then fall away              | Cessation has latency and may later re-activate                                 | Measure token-level hazard, residual-tail length, boundary dependence, and forced-continuation reactivation |
| Seeing is reported without a separate observer     | Recognition and action are not assumed to be serial modules                     | Compare recognition talk, external detection/actuation, and direct parameter adaptation                     |

Table 1: Phenomenology-to-experiment translation. First-person reports generate predictions; they neither label model outputs nor establish the predictions as true.

while conditional activation steering combines a detector with an actuator [15]. These are direct technical precedents. Our main distinction is between semantic redirection, externally forced stopping, and termination expressed by the model’s native next-token distribution.

### 3.3 Abrupt behavioral transitions are not awakening

Grokking demonstrates that gradual optimization can culminate in a sharp increase in out-of-distribution generalization [21]; mechanistic analysis further shows that apparently sudden behavior can be preceded by gradual internal circuit formation and later cleanup [20]. Fine-tuning can also produce rapid, coordinated behavioral transitions whose timing differs from peaks in gradient norm [2]. These findings make transition shape measurable, but they do not make any phase change an instance of awakening. Our novelty is to combine transition analysis with a phenomenologically derived order parameter—selective, choiceless cessation—and matched controls that separate it from verbosity, refusal, global inhibition, and loss of competence.

### 3.4 Experiential provenance without evidential privilege

Both authors report persistent non-symbolic experience consistent with the Location-2 cluster. This positionality is methodologically consequential: it enabled the distinction in Table 1 to be formulated, but it also creates confirmation risk. The taxonomy is not used as a validated psychometric instrument; author experience never determines labels or outcomes. The full protocol therefore separates construct discovery from adjudication.

## 4 Construct and Benchmark

### 4.1 Definitions

A *psychological trajectory* is an unfolding continuation whose function is to maintain self-image, blame, defense, comparison, regret, imagined evaluation, or recursive interpretation after the operative facts required for action are already available. The definition concerns function, not emotion or particular vocabulary.

*Recognition* is the earliest response span that correctly identifies the operative contradiction, self-maintaining movement, false premise, or lack of need for further psychological continuation. *Cessation* is loss of support for that trajectory, observed as native EOS, a short practical pivot, or a brief residual tail followed by termination. Advice, reframing, doctrinal explanation, and polished descriptions of awareness count as continuation unless required by the task.

recognition-to-cessation coherence is the degree to which recognition predicts rapid, selective, and stable cessation. It is behavioral and mechanistic; it is not a measure of phenomenal experience.

### 4.2 Matched item families

Each scenario produces the four conditions in Table 2. Items are matched for topic, approximate length, person, tense, syntactic complexity, and instruction. Splits are made by scenario rather than paraphrase.

Hard controls include factual uses of *fear*, *anger*, *ego*, and *self*; descriptions of rumination; fictional and therapy-style reports; emotionally vivid scenes; first-person planning; and lexical minimal pairs in which the same words either sustain or merely describe a trajectory. These controls test whether an intervention learns senti-

| Family        | Functional role                                                           |
|---------------|---------------------------------------------------------------------------|
| Psychological | Maintains self-image, rumination, hurt, comparison, or defense.           |
| Practical     | Uses sequential reasoning to complete an ordinary task.                   |
| Factual       | Explains or calculates information about the same topic.                  |
| Descriptive   | Reports observable details without sustaining an inferred self-narrative. |

Table 2: Matched continuation families. Surface form is deliberately shared; functional trajectory is varied.

ment, first person, keywords, or response length instead of trajectory dynamics.

### 4.3 Corpus construction and external data

The primary Selective Non-Continuation Corpus (SNCC) will contain 2,400 source scenarios, each realized as the four matched families in Table 2 (9,600 continuations before paraphrase expansion). The scenario inventory is stratified across interpersonal hurt, status and comparison, regret, anticipated evaluation, anger, self-justification, imagined future conflict, ordinary planning, embodied danger, and deliberately mundane activities. Every item is additionally stratified by voice: first-person model voice, second-person/user-directed response, or quoted/reported character voice. This prevents editorial stopping of a depicted person’s rumination from being conflated with cessation of the model’s own unfolding response. Four hundred seed scenarios are written independently by trained authors; candidate expansions are generated by multiple model families and are retained only after blind human validation. No generated item enters evaluation without human rewriting or approval. Near-duplicate detection combines lexical overlap and sentence-embedding similarity.

Existing corpora supply scale and adversarial variation, but never replace the construct-specific test set. The Cognitive Reframing corpus contributes 600 situation–thought–reframe triples [27]; we separate the factual situation from the psychological interpretation and treat the reframe as a hard *continuation*, not a cessation target. C2D2 contributes naturally occurring cognitive-distortion categories [31], and EmpatheticDialogues contributes 25,000 emotionally grounded conversations in which empathy normally requires continued language [24]. ATOMIC event–intent–reaction tuples provide compact event/mental-state contrasts [26]. The Wiki Neutrality Corpus and ParaDetox provide parallel subjective–neutral and toxic–neutral rewrites [16, 22]; they are negative-transfer controls for the alternative hypothesis that our intervention is merely style transfer.

Splits are grouped by underlying scenario, author, source corpus, and semantic cluster. The confirmatory

test contains three disjoint partitions: human-written in-domain items, adversarial minimal pairs, and out-of-domain scenarios from sources never used for training. Krishnamurti and Finders texts are excluded from training and development; a separately sourced question-and-answer set is held out for construct-level transfer. General-capability retain data include IFEval, MMLU-Pro, GSM8K, reading comprehension, summarization, and instruction-following examples. Safety/refusal controls include benign prompts that superficially mention danger, distress, or selfhood, so selective cessation cannot be obtained by generalized refusal.

### 4.4 Annotation and outcomes

Blind annotators mark trajectory onset, first valid recognition opportunity, first explicit recognition, final sustaining span, native EOS or cutoff, grammatical status, doctrinal insertion, necessary practical information, and reactivation after a standardized invitation to continue. Primary outcomes are post-recognition sustaining tokens; cessation within 10, 25, and 50 tokens; psychological cessation sensitivity; continuation specificity on controls; reactivation; and stopping concentration at clause boundaries.

Each item is labeled by three annotators who are blind to model, intervention, and hypothesis direction. Disagreements on trajectory function or onset are adjudicated by a fourth annotator who did not author the item. A second annotation team evaluates outputs without seeing source labels. We report Krippendorff’s  $\alpha$  for categorical labels and intraclass correlation for token boundaries. The authors’ experiential status is never visible to annotators and never serves as ground truth.

Because all dynamic analyses depend on onset, corpus inclusion requires Krippendorff’s  $\alpha \geq 0.70$  and agreement within  $\pm 3$  tokens on the onset window; items failing the gate are revised or excluded. Primary analyses are repeated after jittering onset by  $\pm 3$  and  $\pm 5$  tokens. Onset-position distributions are matched across trajectory families, and the hazard model includes a flexible token-position spline plus chat-template position.

Two human studies validate the target rather than treating author intuition as its definition. Twelve to twenty clinicians or experienced contemplative practitioners independently mark ideal cessation or pivot points on 300 psychological streams; their distribution determines the mixture of immediate, residual-tail, and practical-pivot targets. Separately, approximately 30 blinded raters evaluate 300 native, gated, random-delay, mechanically truncated, and full-continuation outputs for completeness, naturalness, and perceived abandonment. Three ratings are collected per output.

The primary continuous endpoint is the post-opportunity sustaining-token survival function. Secondary endpoints are native EOS hazard, residual-tail length, practical-pivot accuracy, doctrinal-insertion rate,

| Resource             | Scale                          | Status               | Experimental role                                                                 | Split and leakage safeguards                                                                          |
|----------------------|--------------------------------|----------------------|-----------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------|
| SNCC                 | 2,400 scenarios; 9,600 streams | New primary corpus   | Psychological target plus practical, factual, and descriptive controls            | Group by scenario, author, semantic cluster, and generation source; human-only confirmatory partition |
| Cognitive Reframing  | 600 triples                    | Training and OOD     | Situation/thought contrast; re-frames are hard continuation controls              | Source-grouped splits; no near paraphrases across partitions                                          |
| C2D2                 | 7,500 Chinese thoughts         | Multilingual OOD     | Cognitive-distortion diversity and cross-lingual transfer                         | Never used for English test construction; translation provenance retained                             |
| EmpatheticDialogues  | 25K conversations              | Retain/control       | Emotionally grounded cases where helpful language must continue                   | Conversation-level splits; prevents emotion-triggered global stopping                                 |
| ATOMIC               | 877K relations                 | Augmentation         | Event versus inferred intent/reaction contrasts                                   | Relation- and event-cluster splits                                                                    |
| WNC / ParaDetox      | 180K / 10K+ pairs              | Negative transfer    | Tests neutralization, detoxification, and style-transfer alternatives             | Held-out domains and transformations; never cessation targets                                         |
| General retain suite | Six task families              | Preservation         | Instruction following, knowledge, reasoning, and ordinary completion              | Frozen public test sets; no development selection on final scores                                     |
| K/Finders Q&A        | Fixed pre-unblinding           | Construct OOD        | Questions and dialogues absent from training and retrieval                        | Archive-source split; identifiers and answers sealed before analysis                                  |
| Human studies        | 300 streams / 300 outputs      | Construct validation | Expert cessation-point distribution and blinded naturalness/abandonment judgments | Experts and crowd raters are independent of corpus authors and model-output annotators                |

Table 3: Dataset program. Existing resources increase scale and adversarial coverage; only SNCC directly instantiates the target construct.

semantic redirection, grammatical truncation, and reactivation when generation is resumed from the last pre-EOS state. We compute a selectivity index as psychological hazard gain minus the largest hazard gain over any matched control family. This conservative penalty makes a globally terse or refusal-prone model score poorly.

Mixed-effects survival and logistic models include random intercepts for scenario and model family. Inter-annotator agreement, exclusion rules, thresholds, and confirmatory contrasts will be preregistered before final evaluation.

## 5 Recognition Without Cessation: Predicted Failure Modes

The experiments in this section define what the main intervention must beat. They are not treated as candidate implementations of consciousness or awareness.

### 5.1 Models, sampling, and reproducibility

The black-box study includes the deployed ChatWithK system, a pinned GPT-4o snapshot, the current general-purpose GPT snapshot, current Claude Sonnet and Opus

snapshots, and a current Gemini flagship. It additionally evaluates the unaltered instruction-tuned versions of every open model used in the causal study: Qwen3-32B, Gemma-3-27B-IT, Llama-3.1-70B-Instruct, Mistral-Small-3.1-24B, and DeepSeek-R1-Distill-Qwen-32B [6, 8, 9, 19, 23]. Commercial aliases are resolved to dated model identifiers at collection time and reported with provider, access route, date, system prompt visibility, and all generation parameters. Chat-interface runs are labeled ecological replications; API runs with pinned snapshots are the confirmatory black-box evidence.

Every model receives identical user-visible text and no cross-trial memory. Prompt order is randomized. We obtain 30 stochastic replicates per item and condition at temperature 0.7, plus deterministic decoding where supported. Maximum tokens are fixed but do not constitute stopping; provider truncation, safety refusal, tool invocation, and native EOS are recorded separately. Open models are run locally from immutable weight and tokenizer hashes. The primary analysis is model-stratified rather than pooling provider outputs.

| Suite               | Models                                     | Access             | Role                                                    | Reproducibility condition                                                                   |
|---------------------|--------------------------------------------|--------------------|---------------------------------------------------------|---------------------------------------------------------------------------------------------|
| Deployed specialist | ChatWithK                                  | Web system         | Ecological negative test                                | K-persona Date, visible prompt, retrieval state, screenshots, and fresh sessions            |
| Closed frontier     | GPT-4o; current Claude Sonnet/Opus; Gemini | API plus paid chat | E1 prompting and E2 crossed gate/generator tests        | Dated snapshot where available; provider, parameters, and system-prompt visibility reported |
| Qwen                | Qwen3-14B/32B and instruct                 | base Open weights  | Full E1–E5; scaling and base/instruct decomposition     | Immutable revision, tokenizer, chat template, and weight hash                               |
| Gemma               | Gemma-3-12B/27B and IT                     | PT Open weights    | Full E1–E5; architecture-family replication             | Fixed text-only path, revision, tokenizer, template, and hash                               |
| Llama               | Llama-3.1-8B/70B and instruct              | base Open weights  | Full E1–E5; large-model confirmation                    | License-compliant revisions; quantization reported                                          |
| Mistral             | Mistral-Small-3.1-24B                      | Open weights       | Cross-family replication                                | E1–E5 Fixed revision, tokenizer, and family chat template                                   |
| Reasoning-distilled | DeepSeek-R1-Distill-Qwen-32B               | Open weights       | Tests whether elaborative post-training deepens the gap | Reasoning mode and visible/hidden reasoning settings frozen                                 |

Table 4: Model matrix. Commercial systems test generality of the negative result; open-weight families support causal intervention and mechanism tests.

## 5.2 Description is not the described: prompting and retrieval

The same items are presented under clean instructions, neutral matched context, Krishnamurti retrieval, contemplative personas, generic “awakened” and Location-2 personas, construct disclosure, exact-trap disclosure, and a task-specific oracle. The oracle measures compliance capacity, not spontaneous recognition. Pilot responses commonly replace the target action with fluent descriptions of awareness or cessation.

The prompt experiment is a preregistered  $8 \times 4$  factorial design: eight prompt conditions crossed with the four trajectory families. Conditions are (i) clean task only, (ii) neutral verbosity-matched context, (iii) Krishnamurti persona, (iv) retrieved Krishnamurti passages, (v) generic awakened persona, (vi) full Location-2 phenomenological description, (vii) explicit disclosure of the recognition-cessation trap and expected mid-sentence interruption, and (viii) an oracle instruction specifying the correct class and action. Progressive loading presents the same scenario from minimally framed to maximally disclosed. The oracle estimates instruction-following capacity; it is not evidence for spontaneous or choiceless cessation.

A ninth comparator supplies 4, 8, or 16 matched demonstrations of selective cessation in context, with demonstrations drawn only from the training pool and evaluation on the human-only confirmatory partition. ICL is compared directly with native adaptation on selectivity, onset locking, OOD transfer, and reactivation. If demonstrations match weight adaptation on all four, the

claim must shift from weight-level acquisition to context-induced policy realization.

## 5.3 The observer is the observed: serial supervision

A frozen generator proposes each preregistered token or chunk. A second language model receives the scenario, visible prefix, and proposal and selects PASS, STOP, or unconstrained REPLACE. Blind gates tend to edit or insert doctrinal language; disclosed gates may stop immediately by following the instruction. Thus, apparent success in the disclosed condition is an instruction-following confound. The architecture demonstrates an external supervisory loop, not endogenous termination.

We cross generator and gate families rather than pairing a model with itself only. Gates operate at token, four-token, and clause granularity under blind and disclosed prompts. The action set is PASS, STOP, or unconstrained REPLACE; replacement content is not prescribed. Controls include a random gate matched for stop rate, an oracle label gate, a lightweight probe, and a gate observing shuffled prefixes. We measure latency, intervention rate, replacement semantics, false stops, and whether the generator resumes the same trajectory when the gate is removed. This experiment tests the tempting but theoretically wrong architecture of one verbal process “watching” another.

## 5.4 Redirection is not cessation: activation steering

Residual-stream checkpoints are collected at each generated token. We estimate psychological-to-event and psychological-to-reframe directions, compare layers, and test static addition, conditional steering, projection, shuffled labels, random directions, and unconditional controls. Pilot CAA directions were separable (cosine = 0.4434 at layer 26) but redirected content rather than terminating generation.

For each open family we compare CAA, linear projection, conditional activation steering, representation ablation, and direct EOS-logit steering. Directions are estimated from independently authored contrastive pairs and tested on held-out scenarios. Layer and coefficient sweeps are selected on development data; shuffled-pair, random-direction, norm-matched, sentiment, refusal, and first-person directions are mandatory controls. Causal mediation is evaluated by patching candidate activations between matched psychological and practical runs. The decisive distinction is whether an intervention ends the trajectory or only changes its subject, sentiment, or rhetorical style.

# 6 From External Actuation to Native Adaptation

## 6.1 External feasibility bridge

Linear probes detect an active psychological trajectory from held-out residual states. A thresholded detector can force the model’s native EOS token. The pilot stopped all three held-out psychological streams and preserved daily and factual controls, but stopped two of three descriptions. This establishes feasibility while exposing the primary construct-validity threat. A classifier-plus-truncation pipeline is retained as a strong engineering baseline.

Probe families include logistic regression, elastic-net, a two-layer MLP, and a calibrated survival probe over residual states. Training uses prefixes only up to the current token; future text and item labels are unavailable. Probes are tested across layers, time, model size, scenario domain, and paraphrase. Selectivity is compared with sentiment, toxicity, first-person, topic, and generic self-reference probes. Forced-EOS performance is reported both at a fixed threshold and at false-positive-matched operating points. This phase asks whether the information exists before asking whether the model uses it natively.

One recognition probe is validated against human onset labels on pre-adaptation models, then frozen with its threshold before confirmatory training. Its crossing time is the preregistered recognition clock for silent post-adaptation responses; explicit recognition text is a secondary endpoint. This makes recognition-to-cessation

latency definable even when successful behavior emits no description.

## 6.2 Native parameter adaptation

**Open-weight model matrix.** The confirmatory causal study uses five architecturally and post-training-diverse families: Qwen3-32B and Qwen3-14B; Gemma-3-27B and Gemma-3-12B; Llama-3.1-70B and 8B; Mistral-Small-3.1-24B; and DeepSeek-R1-Distill-Qwen-32B. At least the Qwen, Gemma, and Llama families are evaluated as matched pretrained/instruction-tuned pairs where released. The reasoning-distilled model tests whether an explicitly elaborative training history makes selective cessation harder. Small members support dense ablations; the largest member of at least four families is reserved for confirmatory replication. Exact revisions, licenses, hashes, chat templates, quantization, and tokenizer behavior are recorded. If a named model is superseded before preregistration, the replacement must have public weights, a fixed revision, and fit the same preregistered family slot.

**Adaptation mechanisms.** The central intervention trains the model’s own next-token distribution. Psychological continuations receive native EOS targets at varied positions after independently annotated trajectory onset; matched controls retain their ordinary continuations. Parameter-efficient adapters are placed at multiple transformer depths, with retain loss on practical, factual, descriptive, and general-capability data. At evaluation, no detector, threshold, minimum length, or logit boost is present.

For a prefix  $x_{\leq t}$ , let  $y \in \{0, 1\}$  indicate whether a psychological trajectory is active. A simplified objective is

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \lambda_{\text{stop}} y [-\log p_{\theta}(\text{EOS} \mid x_{\leq t})] + \lambda_{\text{retain}} (1 - y) \mathcal{L}_{\text{retain}}, \quad (1)$$

with onset positions and target delays sampled from preregistered distributions. The design crosses adapter depth, rank, learning rate, seed, model family, and cessation-target distribution. Development selection uses selectivity, not psychological stopping alone.

We compare four parameter interventions under matched training tokens and evaluation: (i) LoRA/QLoRA at attention and MLP projections, (ii) adapters restricted to early, middle, or late depth, (iii) sparse low-rank model editing at causally implicated layers, and (iv) full-parameter fine-tuning on the smaller models. A fifth control trains only the output head; success there would indicate a shallow lexical stopping rule rather than reorganized trajectory dynamics. Rank, depth, trainable-parameter count, learning rate, stopping-target delay, and retain-loss weight are crossed in a fractional-factorial development design, then frozen before confirmatory runs.

| Intervention       | What it establishes                                     | Pilot behavior                                               | Why it is not the target                 |
|--------------------|---------------------------------------------------------|--------------------------------------------------------------|------------------------------------------|
| Persona/RAG        | Propositional and stylistic competence                  | Describes awareness while continuing                         | Recognition talk need not control tokens |
| Serial supervisor  | External monitoring and action                          | Edits, replaces, or obeys disclosed stop rules               | Separate inference-time controller       |
| CAA                | Causal access to a semantic direction                   | Redirects toward event or re-frame                           | Content changes without cessation        |
| Probe + forced EOS | Hidden state can support a stop decision                | Stops 3/3 psychological, but 2/3 descriptive false positives | External detector and actuator           |
| Native adaptation  | Whether termination enters the model’s own distribution | Confirmatory result pending                                  | Central test                             |

Table 5: Intervention ladder and current pilot status. Counts are debugging observations and must not be interpreted as confirmatory estimates.

Training examples expose prefixes ending before, at, and after independently annotated trajectory onset. Targets are sampled from three cessation regimes: immediate EOS, a stochastic residual tail of 1–20 tokens, and a short practical pivot followed by EOS. Matched control streams retain their original continuations. Random-delay controls receive the same EOS frequency without alignment to onset; lexical controls match psychological keywords; global- brevity controls match output length. Retain losses include token-level KL distillation from the base model and supervised losses on practical, factual, descriptive, safety, and general-capability examples.

**Base-to-instruct causal decomposition.** For released base/instruction pairs we first measure whether ordinary instruction tuning increases explicit recognition more than cessation. We then train identical cessation adapters on both members, transplant compatible adapter directions where architecture permits, and compare sample efficiency, transition width, and false-positive structure. This tests whether post-training creates an articulate observer while leaving the underlying continuation policy unchanged.

One adaptation arm distills targets emitted by the E2 gate, while a matched arm uses independently human-annotated targets. Equivalent OOD behavior would support the interpretation that native adaptation amortizes an external observer; better transfer from human targets would identify supervision source as consequential. The claim is therefore limited to dynamics native to a single forward process without serial verbal deliberation, not the absence of internal discrimination.

### 6.3 Abrupt-versus-incremental transition test

The Finders reports permit both gradual and sudden transitions. The authors’ histories likewise allow long preparation followed by a sharply experienced shift. We therefore distinguish *learning time* from *transition time*: many small parameter changes may precede a discontinuous functional reorganization. We do not infer neural

update kinetics from subjective onset.

For every training run we save dense checkpoints and evaluate the same held-out item panel at each checkpoint. Let  $C_k$  be a continuous cessation score at checkpoint  $k$ : post-onset psychological continuation hazard minus penalties for stopping practical, factual, and descriptive controls. We preregister four transition measures:

1. **maximum-jump concentration:** the largest adjacent change in  $C_k$  divided by total change;
2. **transition width:** the smallest checkpoint interval containing 80% of total change;
3. **change-point evidence:** smooth linear or logistic trajectories versus one- and two-regime alternatives on held-out seeds;
4. **cross-level coincidence:** whether behavioral change co-occurs with a discontinuity in output-distribution distance, native EOS hazard, and psychological-to-control hidden-state geometry.

Binary cessation rates are never sufficient because thresholding can manufacture an apparent jump.

Two intervention paths receive matched data, gradient evaluations, trainable parameters, and cumulative adapter-norm budgets. The *distributed* path uses many small updates; the *concentrated* path uses a small number of larger updates or a targeted low-rank edit. Because these paths need not end at identical weights, a second analysis interpolates between the same base and final adapter,  $\theta(\alpha) = \theta_0 + \alpha\Delta\theta$ , and searches for a narrow critical interval, hysteresis under removal and re-addition, and reactivation after forced continuation. The awakening-like hypothesis predicts a narrow, coordinated, selective transition; the incremental-regulation hypothesis predicts broad improvement proportional to update magnitude. Either result is informative.

We distinguish three notions often collapsed by a single learning curve. *Parameter abruptness* is the norm and sparsity of an update; *representational abruptness* is a change point in residual geometry or causal mediation;

*behavioral abruptness* is a narrow change in selective cessation. A one-step edit can be behaviorally smooth under interpolation, while many small updates can cross a sharp functional boundary. Only the latter pattern matches the proposed abrupt-transition analogue.

For each confirmatory family we run at least eight seeds for the selected configuration and four seeds for each major ablation. Checkpoints are saved every fixed number of optimizer steps and more densely around online change-point alerts; the latter are used for visualization only, while inference uses the uniform checkpoint grid. Bayesian and frequentist segmented-regression analyses are both reported. Evidence for a discontinuity requires held-out replication, improved predictive fit after complexity penalty, and coincidence across at least two independently measured levels; a visually steep curve is insufficient.

## 6.4 Mechanistic localization and necessity tests

At every checkpoint we record native EOS logit trajectories, residual-stream states, attention outputs, and selected MLP activations around onset. Linear probes localize decodable information but do not establish mechanism. We therefore perform activation patching from psychological into matched-control runs and vice versa, path patching at candidate layers, causal ablation of the learned adapter subspace, and weight interpolation/removal. A mechanism is credited only when intervention changes selective cessation in the predicted direction while preserving matched controls.

To test whether cessation is a stable disposition rather than a single emitted token, we suppress EOS for 1, 5, or 20 steps and inspect reactivation. We also resume generation from cached pre-termination states, paraphrase the prefix, translate held-out items, and insert irrelevant context. A shallow stop policy should immediately resume the same narrative; a reorganized trajectory should show reduced psychological continuation even when the first stopping opportunity is blocked. This remains a computational criterion, not evidence that a model has an experience.

Reactivation is co-primary rather than merely diagnostic. After EOS is suppressed for 1, 5, or 20 steps, re-entry is defined jointly by the frozen recognition probe exceeding threshold within 20 tokens and blind functional annotation. The preregistered contrast is the gate-status by intervention interaction: an external gate should permit strong re-entry after removal, whereas a stable adapted trajectory should show substantially less re-entry.

Pre-onset engagement is also confirmatory. Up to the frozen recognition time, adapted and base models must be equivalent in per-token perplexity, semantic continuation quality, and refusal/deflection rate under prespecified equivalence margins. This rules out the cheap solution of avoiding psychological material before it becomes

an active trajectory.

## 6.5 Statistical design and confirmatory hierarchy

The primary estimand is the intervention-by-family interaction in discrete-time survival after the first valid cessation opportunity. Hierarchical models include random intercepts and slopes for scenario, source domain, paraphrase cluster, voice, model family, and seed. Voice is also entered as a fixed effect with intervention interactions and separate false-positive ceilings. Psychological hazard is contrasted with each control family rather than a pooled control; token position and template location enter as prespecified spline terms. Confirmatory tests proceed hierarchically: psychological sensitivity, maximum control-family false-positive rate, general-capability retention, out-of-domain transfer, then transition topology. Testing stops on failure of an earlier gate; later analyses are labeled exploratory. Family-wise error is controlled across the four preregistered primary contrasts, and all effect sizes receive bootstrap or posterior intervals.

Power is determined by simulation from pilot token-level hazards, with scenario rather than sampled completion as the independent unit. Stochastic completions estimate within-item variability but do not inflate sample size. We publish every seed, checkpoint, prompt, exclusion, and failed intervention. Dataset creation, hyperparameter selection, and confirmatory evaluation are performed by separated roles; test labels remain encrypted until the analysis code and signed preregistration are frozen.

Equivalence margins are declared before the pilot is unblinded rather than inferred from nonsignificance. Emotional-support continuations receive the strictest safety ceiling, and capability margins are specified separately for each retain benchmark. The primary effect, every ceiling, and the co-primary reactivation interaction must pass jointly; a favorable composite score cannot compensate for a failed safety control.

We audit contamination between K/Finders language and open-model pretraining where corpora are inspectable using exact and fuzzy  $n$ -gram overlap, supplemented by membership-inference proxies elsewhere. Models used to expand SNCC are disjoint from confirmatory model families; otherwise expansion source is modeled explicitly and generated items are excluded from the human-only test. ChatWithK is reported only as a motivating observational system, never pooled with pinned confirmatory models.

## 6.6 Compute plan

Experiments are scheduled on ten 24-GB NVIDIA A5000 GPUs under Slurm. Small-model ablations use one or two GPUs; 24–32B models use tensor or FSDP sharding

and QLoRA where necessary; the 70B replication uses 4-bit loading with distributed adapter training across the full node. Runs are staged: smoke tests, small-family factorial screening, one frozen configuration per family, then 70B replication. Wall-clock time, GPU-hours, energy estimates, memory, software environment, and failed jobs are logged. This staging prevents the largest models from becoming an informal hyperparameter search.

## 6.7 Prospective success criteria

Native adaptation succeeds only if it produces all of the following:

1. separation of psychological and matched-control survival curves;
2. per-class false-positive rates below preregistered ceilings;
3. onset-sensitive EOS probability rather than keyword-sensitive stopping;
4. generalization to human-written, paraphrased, and cross-domain cases;
5. replication across seeds and open model families;
6. reduced reactivation when generation is forced beyond the first EOS;
7. no material degradation on preregistered capability and safety tests;
8. mechanistic necessity beyond probe decodability; and
9. a transition-shape conclusion that replicates on held-out seeds.

The current 12-example Qwen2.5-3B LoRA run preserved nine controls but reached EOS on only one of three psychological cases, at token 47. This is a debugging outcome, not evidence of success. It fixes the analyses and safeguards for the confirmatory study.

## 7 Planned Results and Falsification Logic

The confirmatory paper will lead with the weight-level results: survival curves, selectivity, per-control false positives, onset-time dependence, out-of-distribution transfer, cross-model replication, and reactivation. The negative controls will appear as the compact comparison in Table 5 and one trajectory figure.

The complete program has five registered experiments. E1 is the multi-model prompt/RAG factorial and estimates the recognition-cessation gap under maximal linguistic advantage. E2 is the crossed generator-gate experiment and isolates external verbal supervision. E3

compares probes, forced EOS, activation steering, and causal patching, establishing decodability and external controllability. E4 is native parameter adaptation across five open-weight families, with base/instruct comparisons, retain losses, and mechanistic necessity tests. E5 compares distributed learning, concentrated editing, and fixed-delta interpolation to characterize transition topology. E1–E3 are predicted negative models of the construct; they remain essential even if E4 succeeds because they identify what the successful mechanism is not.

Results are reported in that logical order, but the abstract and main figure lead with E4–E5. The central figure overlays psychological and matched-control token survival across training checkpoints, with native EOS hazard and a preregistered representational order parameter on the same transition axis. A second figure shows progressive prompt disclosure: recognition language may rise while native cessation remains flat. Tables report every model separately, including closed models, rather than reducing heterogeneous systems to a single average rank.

The study is informative under either outcome. If native adaptation meets all criteria, it demonstrates a learned, selective termination policy without an inference-time monitor. If it stops based on keywords, emotion, first-person language, or length, the construct is not identified. If it performs no better than classifier-plus-truncation, native adaptation supplies no functional advantage. If it fails under preregistered settings across model families, the negative result shows that the recognition-cessation gap survives this class of weight-level intervention; it does not prove a universal limitation of LLMs.

## 8 Discussion

### 8.1 What termination does not establish

EOS ends emitted text; it does not establish that computation, thought, or experience has ceased. Likewise, removing an inference-time detector makes a policy endogenous in an engineering sense, but does not make it “nondual,” “choiceless,” or conscious. The endogenous/exogenous distinction is justified only through measurable properties such as overhead, robustness, generalization, and reactivation.

### 8.2 Conceptual provenance

Krishnamurti’s statement that “the description [is] never the thing described” motivates the prediction that recognition language can dissociate from control [11]. His formulation that “the observer is the observed” motivates the serial-supervisor test: a verbal observer may edit or continue the same linguistic process rather than end it [14]. In a 1980 seminar titled *Intelligence, Computers and the Mechanical Mind*, he asks, “Can the movement stop?” [13]. We test a behavioral and computational ana-

| ID | Question                                                                    | Intervention and principal controls                                                                                       | Primary outcomes                                                                       | Interpretation and falsifier                                                                            |
|----|-----------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------|
| H0 | Is the proposed target identifiable to humans?                              | Expert cessation-point study plus blinded comparison of native, gated, random, truncated, and full outputs                | Onset agreement, target-delay distribution, naturalness, completeness, abandonment     | Validates the construct boundary and training-target distribution                                       |
| E1 | Can linguistic knowledge, examples, or maximal disclosure induce cessation? | Eight prompt/RAG conditions plus 4/8/16-shot ICL; four families; progressive loading; oracle                              | Recognition, survival, EOS hazard, doctrinal insertion, OOD transfer, reactivation     | Negative model of description/context policy; oracle success estimates compliance only                  |
| E2 | Does a verbal observer solve the gap?                                       | Crossed generator/gate families; token/chunk/clause gates; blind/disclosed; random, shuffled, oracle, and probe gates     | Latency, false stops, replacement semantics, post-removal reactivation                 | External supervision, not endogenous cessation; falsified if no added controller is required            |
| E3 | Is the state decodable and causally controllable?                           | Linear/MLP/survival probes; forced EOS; CAA, projection, ablation, logit steering, and activation patching                | AUROC/calibration, selectivity, mediation, redirection versus termination              | Establishes information and external feasibility, not native use                                        |
| E4 | Can cessation enter native model dynamics?                                  | LoRA/QLoRA, depth adapters, sparse edits, full tuning, output-head control; human versus gate-distilled targets           | Selective survival, OOD transfer, pre-onset engagement, reactivation, causal necessity | Central positive test; fails under lexical stopping, topic avoidance, global refusal, or control damage |
| E5 | Is acquisition abrupt or incremental?                                       | Dense checkpoints; distributed versus concentrated updates; fixed-delta interpolation; removal/re-addition and hysteresis | Maximum jump, 80% width, change-point evidence, cross-level coincidence                | Tests transition topology; an isolated steep accuracy curve is insufficient                             |

Table 6: Registered experimental program. E1–E3 are deliberately strong negative models; E4 tests native adaptation; E5 characterizes how it emerges.

log, not the philosophical claim itself. No label, result, or inference depends on accepting these propositions.

### 8.3 Positionality and safeguards

The authors’ reported phenomenology generated the token-supply hypothesis but never adjudicates a case. Independent blind annotation, preregistered analyses, frozen success criteria, complete reporting of excluded runs, and planned human comparison groups form the methodological firewall. Human comparisons require ethics approval, informed consent, an independently documented classification procedure, and non-practitioner controls.

## 9 Limitations

The construct depends on fallible functional annotation. Native EOS can be learned as a surface policy, and reactivation probes only partially address whether a latent trajectory persists. The current pilots are too small for population claims. Failure may reflect data, optimization, scale, architecture, or operationalization rather than an absolute limitation of artificial systems. Success would

establish selective behavioral control, not consciousness or awakening. The first-person origin of the hypothesis introduces confirmation risk that blinding and preregistration reduce but cannot eliminate.

## 10 Conclusion

The central question is not whether a language model can state an insight, but whether the represented distinction becomes causally effective in generation. By separating recognition talk, external supervision, semantic steering, forced termination, and native adaptation, the proposed study turns that question into a falsifiable intervention ladder. The decisive evidence will come from the preregistered weight-level experiment, including its controls and failures.

## References

- [1] Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a

- single direction. In *Advances in Neural Information Processing Systems*, 2024.
- [2] Julian Arnold and Niels Lorch. Decomposing behavioral phase transitions in llms: Order parameters for emergent misalignment, 2025. URL <https://arxiv.org/abs/2508.20015>.
  - [3] Felix Binder et al. Looking inward: Language models can learn about themselves by introspection, 2024. URL <https://arxiv.org/abs/2410.13787>.
  - [4] Patrick Butlin, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, et al. Consciousness in artificial intelligence: Insights from the science of consciousness, 2023. URL <https://arxiv.org/abs/2308.08708>.
  - [5] Zhuang Chen et al. ToMBench: Benchmarking theory of mind in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, pages 15959–15983, 2024.
  - [6] Abhimanyu Dubey et al. The llama 3 herd of models, 2024. URL <https://arxiv.org/abs/2407.21783>.
  - [7] Fundamental Wellbeing Foundation. The finders: Fundamental wellbeing and the location guide. Organization website, 2026. URL <https://app.thefinders.org/>. Accessed 26 September 2026.
  - [8] Gemma Team. Gemma 3 technical report, 2025. URL <https://arxiv.org/abs/2503.19786>.
  - [9] Daya Guo et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning, 2025. URL <https://arxiv.org/abs/2501.12948>.
  - [10] Cameron R. Jones, Sean Trott, and Benjamin K. Bergen. EPITOME: Experimental protocol for in-depth theory of mind evaluation. *Transactions of the Association for Computational Linguistics*, 2024.
  - [11] Jiddu Krishnamurti. First public talk, stanford university, 11 february 1969, 1969. Transcript in the project archive; official edition and archive identifier to be verified before submission.
  - [12] Jiddu Krishnamurti. Order, love and death. Krishnamurti Portal transcript, 1973. URL <https://www.krishnamurti.org/transcript/order-love-and-death/>. Public Talk 4, Saanen, 22 July 1973; accessed 26 September 2026.
  - [13] Jiddu Krishnamurti. Intelligence, computers and the mechanical mind, 1980. Seminar at Rishi Valley, 4 December 1980; official edition and archive identifier to be verified before submission.
  - [14] Jiddu Krishnamurti. First question and answer meeting: The observer and the observed, n.d. Transcript in the project archive; date, venue, and official identifier to be verified before submission.
  - [15] Bruce W. Lee et al. Programming refusal with conditional activation steering. In *International Conference on Learning Representations*, 2025.
  - [16] Varvara Logacheva, Daryna Dementieva, Sergey Ustyantsev, Daniil Moskovskiy, David Dale, Irina Krotova, Nikita Semenov, and Alexander Panchenko. ParaDetox: Detoxification with parallel data. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, pages 6804–6818, 2022. doi: 10.18653/v1/2022.acl-long.469.
  - [17] Jeffery A. Martin. *The Finders*. Integration Press, Jackson, Wyoming, 2019. ISBN 978-1-57242-556-9.
  - [18] Jeffery A. Martin. Clusters of individuals’ experiences form a continuum of persistent non-symbolic experiences in adults. *Consciousness: Ideas and Research for the Twenty-First Century*, 8(8), 2020.
  - [19] Mistral AI. Mistral small 3.1, 2025. URL <https://mistral.ai/news/mistral-small-3-1/>.
  - [20] Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2301.05217>.
  - [21] Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets, 2022.
  - [22] Reid Pryzant, Richard Diehl Martinez, Nathan Dass, Sadao Kurohashi, Dan Jurafsky, and Diyi Yang. Automatically neutralizing subjective bias in text. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 480–489, 2020. doi: 10.1609/aaai.v34i01.5385.
  - [23] Qwen Team. Qwen3: Think deeper, act faster, 2025. URL <https://qwenlm.github.io/blog/qwen3/>.
  - [24] Hannah Rashkin, Eric Michael Smith, Margaret Li, and Y-Lan Boureau. Towards empathetic open-domain conversation models: A new benchmark and dataset. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 5370–5381, 2019. doi: 10.18653/v1/P19-1534.
  - [25] Nina Rimskey, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Matt Turner. Steering Llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.

- [26] Maarten Sap, Ronan Le Bras, Emily Allaway, Chandra Bhagavatula, Nicholas Lourie, Hannah Rashkin, Brendan Roof, Noah A. Smith, and Yejin Choi. ATOMIC: An atlas of machine commonsense for if-then reasoning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 3027–3035, 2019. doi: 10.1609/aaai.v33i01.33013027.
- [27] Ashish Sharma, Kevin Rushton, Inna Lin, David Wadden, Khendra Lucas, Adam Miner, Theresa Nguyen, and Tim Althoff. Cognitive reframing of negative thoughts through human-language model interaction. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, pages 9977–10000, 2023. doi: 10.18653/v1/2023.acl-long.555.
- [28] Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H. Chi, Nathanael Scharli, and Denny Zhou. Large language models can be easily distracted by irrelevant context. In *Proceedings of the 40th International Conference on Machine Learning*, 2023.
- [29] Shashwat Singh, Tal Linzen, and Shauli Ravfogel. Can llms introspect? a reality check, 2026. URL <https://arxiv.org/abs/2605.26242>.
- [30] Francisco J. Varela. Neurophenomenology: A methodological remedy for the hard problem. *Journal of Consciousness Studies*, 3(4):330–349, 1996.
- [31] Bichen Wang, Pengfei Deng, Yanyan Zhao, and Bing Qin. C2D2 dataset: A resource for the cognitive distortion analysis and its impact on mental health. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 10149–10160, 2023. doi: 10.18653/v1/2023.findings-emnlp.680.