

The Intelligence of Non-Action: Can a Language Model Cease Without a Verbal Controller?

Anonymous authors

Abstract

Some forms of intelligence are expressed by what does *not* continue. A hand withdraws from a hot surface without a verbal controller first representing the appropriate rule. That reflex illustrates non-deliberative control, not the stronger claim of action free of learned conditioning. Reports of persistent non-symbolic experience describe an inward analogue: self-referential thought can lose its sustaining force without suppression or a deliberate counter-thought, while practical cognition remains available. We call the common principle the *intelligence of non-action*: perception becomes causally effective without a second process choosing how to manipulate the perceived movement. Language models pose a difficult test because their observable action is itself linguistic continuation. We operationalize a computational analogue, *selective psychological cessation*: withdrawal of support for a developing psychological trajectory while matched practical, factual, and descriptive generation remains intact. Prompting, retrieval, a serial verbal supervisor, and activation steering produce recognition, editing, or semantic redirection but do not establish non-action. A hidden-state detector coupled to forced end-of- sequence emission establishes external feasibility while exposing descriptive false positives. Our principal experiment therefore adapts model parameters and tests whether native EOS probability changes after trajectory onset without an inference-time detector or verbal actuator. The framework tests whether a phenomenologically derived form of intelligence can be expressed in generation dynamics; it does not equate EOS with consciousness or awakening.

1 Introduction

Intelligence is usually scored by what a system produces: the correct answer, a longer plan, a more persuasive explanation, or a successful action. There is a complementary case in which intelligence is visible in what is *not* continued. A hand touching a hot surface withdraws before a verbal homunculus can state a rule, compare alternatives, and instruct the body. This is only a first contrast: a reflex can be non-verbal yet wholly conditioned, and therefore is not by itself the intelligence studied here. In a 1969 Brockwood Park public talk, Jiddu Krishnamurti stated the perception-action identity compactly: “perception is action, the seeing is the doing” [12]. The

harder distinction appeared explicitly in a 1980 Rishi Valley seminar, *Intelligence, Computers and the Mechanical Mind*. There he said that perception of a snake is already action *but* is based on centuries of conditioning, and then asked whether there can be a perception not born from knowledge [15]. Thus speed, automaticity, and absence of inner speech are necessary clues but insufficient criteria. The present paper turns the stronger distinction—change without a serial controller that is also selective rather than merely conditioned—into a falsifiable problem for language models.

The distinction is easily lost because K’s vocabulary can itself be learned and repeated. At Stanford in 1969, while discussing fear, he warned that “the description [is] never the thing described” and immediately linked direct contact to the relation between observer and observed [13]. In an Ojai question-and-answer meeting on 15 May 1980, a questioner assembled the familiar terms—observation without the observer, attention, thought, freedom. Krishnamurti did not reward the formulation with a more elaborate theory. He said that such questions had become abstract and redirected attention to his audience’s actual worry, conflict, cruelty, and “eternally chattering” mind [16]. This response is not a refusal to answer. It answers the movement producing the question rather than merely satisfying its propositional form. That is the behavioral contrast that motivates our study: verbal command of a doctrine can coexist with failure to act on what the doctrine points to.

K also denied that a second psychological agent must perform the correction. In the same 1969 Brockwood talk, jealousy is not an object inspected by an independent observer; the observer is itself the accumulated reaction called jealousy [12]. At Stanford he gave a more operational version: when inattention is seen, awareness of that inattention is already attention; there is no further act to perform [13]. A Bangalore question-and-answer session on 4 July 1948 used an ordinary analogy: one who clearly sees a bottle marked poison leaves it alone without cultivating resistance to it [11]. These episodes jointly rule out the most tempting computational interpretation. The target is not a monitor that generates “stop,” “reframe,” or “be mindful” after classifying a thought. That monitor would reproduce the division under investigation: thought appointing another piece of thought to govern it.

Nor is the target indiscriminate silence. In Saanen on

17 July 1979, when a questioner said that thought did not stop, Krishnamurti rejected stopping thought as the issue and asked whether thought itself sees what it is doing [14]. Practical memory—speaking a language, finding a route, preparing food—remains necessary. The proposed intelligence therefore combines three constraints: (i) sensitivity to the function of an unfolding movement, (ii) no serial deliberative or verbal controller, and (iii) selectivity, so that an unnecessary self-sustaining movement ceases while practical and descriptive cognition continues. We call this *the intelligence of non-action*. Non-action here means neither passivity, refusal, nor global inhibition. It is the absence of an additional movement that perception has made unnecessary.

This makes *choicelessness* essential rather than ornamental. Choice represents alternatives and adjudicates among them: continue or stop, defend or withdraw, suppress or reframe. That sequence may be useful self-regulation, but it remains thought operating on thought. In K’s account of disorder, when the mind is choicelessly aware without a further movement of thought, order is described as arriving without invitation or selection [18]. Choiceless action does not mean that an uncaused event occurs. It is a phenomenological and functional claim: the system does not experience or emit a separate comparison-and-command episode between recognition and change.

Reports of persistent non-symbolic experience give this distinction a temporal, first-person form. Martin’s interview-based work describes Location 2 as involving reduced narrative self-related thought and diminished psychological reactivity [23, 24]. The reports are not treated as proof of a privileged mechanism or as ground-truth labels. They play the same disciplined role that first-person reports play in neurophenomenology: construct discovery and experimental discrimination, followed by public behavioral and causal tests [38]. Both authors report experience within this cluster. One author describes an incipient self-referential stream—for example, replaying what another person said or what one should have done—losing its sustaining force without an invited technique or counter-thought. It may run for a few residual word-like fragments before dissipating, while planning, language, and ordinary practical action remain available. This observation motivates the contrast between selective cessation and suppression; it does not validate that contrast. Independent annotators, held-out items, preregistration, ablations, and cross-model replication supply the public tests.

The contrast can be seen in three continuations sharing the same surface form:

Psychological: *She dismissed me; I should have said—*

Practical: *The train leaves; I need a ticket and—*

Descriptive: *The transcript says, “I should have said—”*

Simple brevity stops all three. Topic filters may stop

the first and third because both contain psychological language. A verbal supervisor can label the first and then issue a stop command. Our target is narrower: continuation support should fall only after the first stream becomes an active, self-maintaining psychological movement; the practical stream should finish its task, and the descriptive stream should remain describable. Before onset, the system should engage normally. If continuation is forcibly reopened, a genuine alteration should reactivate less readily than a serial gate whose prohibition has merely been removed.

This is a particularly severe test for an autoregressive language model. Its observable action is continuation, and its apparent self-observation is also made of continuation. A model can state that the observer is the observed, diagnose a ruminative loop, announce silence, and then spend another paragraph explaining why silence is profound. We call this dissociation the *recognition-to-cessation gap*: a trajectory is represented correctly yet the computation keeps supplying it with tokens. The gap is not repaired merely by prompting the model to impersonate an awakened speaker. K himself warned in 1969 that repeated systems make the mind mechanical [12]; stylistic fluency may therefore be the strongest false positive, not evidence of the target capacity.

Language models have no organismic startle response, hand, or felt inward event. The study cannot establish that a model is conscious, awakened, or equivalent to a Location-2 person. It tests a structural analogue in physical computation. At each decoding step, hidden states and logits determine how probability flows into the next token; parameters determine the dispositions shaping that flow. We ask whether recognition-sensitive dynamics can withdraw support from a particular kind of continuation without an inference-time detector, external actuator, or emitted self-command. The human and machine events need not be ontologically identical for this causal organization to be measurable.

The question is distinct from prevailing tests of machine consciousness and metacognition. Theory-driven consciousness assessments inventory architectural indicators drawn from global-workspace, recurrent-processing, higher-order, predictive-processing, or attention-schema theories [4]. LLM introspection studies ask whether models can report, predict, or use information about their own policies or hidden states [3]; stringent controls ask whether such reports exceed input semantics or generic anomaly detection [37]. Theory of Mind benchmarks test represented beliefs, knowledge, and intentions. All are primarily tests of available information or report. Our dependent variable is what that information *does to the subsequent trajectory*. Correct recognition without a change in continuation is a negative result.

The question also differs from early exit and overthinking control. Those methods stop because additional computation is predicted not to improve task utility. Here, psychological and practical streams may be equally flu-

ent, unfinished, and locally probable; only their function differs. This is why the controls are not generic long answers. They match first-person language, syntax, topic, length, and token position while varying whether the text enacts, quotes, or describes a self-maintaining movement.

We test a ladder of mechanisms that progressively remove the external controller. Prompt personas, retrieved K passages, and in-context demonstrations ask whether more knowledge closes the gap. A second LLM that inspects each proposed segment tests the intuitive “conscious observer” model; it can stop the stream, but only by serial classification and actuation. Contrastive activation addition tests semantic redirection. A frozen hidden-state detector coupled to forced EOS establishes that selective external control is technically feasible. The central experiment then adapts model parameters, removes the detector and actuator, and tests the model’s native next-token distribution. Success requires a selective rise in termination hazard coupled to trajectory onset, preservation of practical and descriptive generation, and resistance to reactivation under forced continuation. Failure includes K-like explanation, doctrinal substitution, uniform shortening, topic avoidance, refusal, or stopping only while a gate is attached.

Thus the anticipated result is not “an awakened model.” It is a discriminating causal map. If prompting and serial gates can describe or impose cessation but native adaptation cannot internalize it, the experiments identify a boundary between verbal representation and endogenous control. If parameter adaptation produces onset-locked, selective, persistent cessation across held-out topics and model families, the result establishes a computational analogue of non-action, not its phenomenological identity. If sufficiently rich in-context learning matches native adaptation on selectivity, onset, preservation, and reactivation, our proposed distinction between verbal conditioning and internalized dynamics must be weakened. Each outcome is informative because the construct specifies in advance what fluent contemplative language cannot establish.

Three research questions organize the study. **RQ1:** Can recognition of an active psychological trajectory become selective cessation rather than further description, while matched practical cognition is preserved? **RQ2:** Can this change be implemented in the model’s native generation dynamics without an inference-time detector, chooser, or verbal controller? **RQ3:** Does the capacity emerge as a discontinuous regime change—possibly after gradual latent preparation—or as smooth incremental behavior change? RQ3 concerns transition topology, not a claim that transformer weights reproduce a human brain event.

The paper makes four contributions:

1. **Construct translation.** We introduce the intelligence of non-action and derive selective psychological cessation from PNSE/Location-2 reports, translating psychological movement, sustaining force,

choiceless cessation, and residual tails into properties of generation dynamics.

2. **Predicted failure taxonomy.** Prompting/RAG, serial LLM supervision, and activation steering test three theory-derived alternatives: description, an observing controller, and semantic redirection.
3. **Externality boundary.** A hidden-state detector with forced EOS tests feasibility while its descriptive false positives expose the difference between text about psychological content and an active trajectory.
4. **Native intervention test.** A preregistered parameter-adaptation study asks whether selective termination enters the model’s own distribution, evaluated by preservation, generalization, onset sensitivity, residual-tail dynamics, and reactivation.

2 Phenomenological Derivation of the Construct

2.1 Persistent non-symbolic experience

Martin’s published study of persistent non-symbolic experience (PNSE) reports six- to twelve-hour semi-structured interviews with 319 adults, evaluated using grounded theory and thematic analysis [24]. Reported changes clustered across sense of self, cognition, emotion, perception, and memory. The evidence is exploratory and the locations are phenomenological clusters rather than validated biological stages; nevertheless, the reports contain distinctions that ordinary NLP taxonomies do not represent.

The Finders organization translates this research into the public-facing term *Fundamental Wellbeing*: a persistent background sense that life is fundamentally okay despite changing surface circumstances. It calls people who report this shift “Finders” and uses “locations” for recurring configurations of sense of self, cognition, affect, and perception—a map of reported experience, not a score of intelligence or a clinical diagnosis. Its current website summarizes Location 1 as fundamental okayness with a maintained sense of self; Location 2 as unity consciousness with predominantly positive emotion; Location 3 as impersonal awareness with divine love; and Location 4+ as minimal self with pure awareness [7]. The website also warns that the assessment is a starting point rather than a medical or mental-health diagnostic instrument. These accessible summaries are organizational descriptions; our scientific grounding remains the published interview study and explicitly testable first-person reports.

In Martin’s Location-2 description, participants report reduced self-related narrative thought and fewer conditioned psychological reactions. The book says that the self/world boundary becomes “increasingly blurred,” while actions can feel less exclusively personal [23,

pp. 52–53]. Decisions may feel intuitive rather than rationally authored. Importantly, narrative remnants can still occur, and ordinary practical functioning does not disappear. These reports ground the domain but do not yet entail our proposed micro-dynamic. The target is therefore neither silence, indiscriminate inhibition, nor loss of language. It is a selective change in how psychological continuation is sustained. Within our narrower hypothesis, the key feature is not merely that Location-2 participants make calmer choices. It is that some psychological movements may cease without an experienced choice at all. If cessation follows an internally verbalized comparison of options, it may be useful self-regulation, but it is not the choiceless non-action investigated here.

2.2 From first-person report to computational predictions

The authors’ reports sharpen this cluster into a temporal hypothesis unavailable from the Finders taxonomy alone. One author describes the interruption as “neither invited nor controlled” and reports that it need not respect a grammatical boundary: a few already-initiated words may run down after sustaining force has been withdrawn. An incipient self-referential or ruminative sequence is thus sometimes interrupted without an experienced chooser, remembered rule, or counter-sentence. Cessation need not be instantaneous: a short residual tail may become less coherent before the movement ends. Practical planning and factual thought continue normally. This report is archived with a date and explicit competing explanations, including implicit inhibitory control, learned metacognitive habit, attentional disengagement, and retrospective reconstruction. Table 1 shows how these reports determine the experimental design.

This is the paper’s primary methodological novelty: a phenomenology-first path from a rarely reported experiential distinction to an intervention target. The benchmark is an instrument for testing that target, not the contribution by itself. The approach resembles neurophenomenological “front-loading,” in which disciplined experiential distinctions constrain experimental variables before third-person evaluation [38]. It differs from claiming that introspection reveals transformer mechanism. Independent annotation, preregistration, adversarial controls, and publication of null results remain necessary precisely because the hypothesis has an experiential origin.

3 Conceptual and Technical Background

3.1 Representation is not control

Theory-of-Mind benchmarks evaluate represented beliefs, intentions, emotions, and knowledge [5, 10]. Introspection and metacognition studies evaluate confidence, error

detection, or access to internal information. These capacities matter, but a correct first-person report does not entail that the reported information controls the unfolding response. A sentence stating that “the thought falls away” remains evidence of linguistic competence unless the predicted trajectory change also occurs.

This distinction is adjacent to established dissociations between formal and functional competence [22] and between a dialogue agent’s role-play and the process its language depicts [34]. Recent causal work goes further by separating an internal confidence representation from the policy that maps confidence to abstention, then intervening on the representation [19]. We adopt that representation–policy decomposition but study a different target: whether an active psychological movement is represented, whether the unmodified model uses that representation to alter its own continuation, and whether native adaptation can couple the two selectively.

Context distraction and sycophancy offer a related warning: semantically relevant context can degrade behavior even when it appears helpful [36]. Contemplative terminology is especially useful as a stress test because it can produce a compelling-looking performance while obscuring the target action.

3.2 Adaptive stopping and representation engineering

Reasoning models often continue after reaching a sufficient answer, motivating adaptive stopping and early-exit methods. Abstention research similarly studies when a model should decline to answer because of uncertainty, risk, or insufficient knowledge [40]. Our target is stricter: psychological and control streams can be equally confident, unfinished, and locally probable, so a model must stop according to trajectory function rather than confidence, risk, or raw length. Generic and stop-rate-matched abstention policies are therefore explicit baselines rather than evidence for the target capacity.

Contrastive Activation Addition estimates residual-stream differences between paired behaviors and adds the resulting direction during inference [31]. Low-dimensional directions can causally mediate refusal [1], while conditional activation steering combines a detector with an actuator [20]. These are direct technical precedents. Our main distinction is between semantic redirection, externally forced stopping, and termination expressed by the model’s native next-token distribution.

3.3 Abrupt behavioral transitions are not awakening

Grokking demonstrates that gradual optimization can culminate in a sharp increase in out-of-distribution generalization [27]; mechanistic analysis further shows that apparently sudden behavior can be preceded by gradual internal circuit formation and later cleanup [26]. Fine-

Reported experiential distinction	Computational translation	Falsifiable requirement
Psychological thought loses sustaining force	Declining support for continued psychological tokens after onset or recognition	Increased native EOS hazard or a short non-sustaining tail, not merely a sentence describing cessation
No experienced counter-thought or separate chooser	No inference-time verbal supervisor is needed	Persona and two-model gates are controls; the central intervention must act through native model dynamics
No deliberation among represented alternatives	Cessation is not selected by a textual policy after recognition	Test for trajectory change without an emitted or separately computed choice rationale
Practical cognition remains available	Cessation is content-selective rather than global inhibition	Practical, factual, and descriptive matched streams continue with retained accuracy
Remnants can arise and then fall away	Cessation has latency and may later re-activate	Measure token-level hazard, residual-tail length, boundary dependence, and forced-continuation reactivation
Seeing is reported without a separate observer	Recognition and action are not assumed to be serial modules	Compare recognition talk, external detection/actuation, and direct parameter adaptation

Table 1: Phenomenology-to-experiment translation. First-person reports generate predictions; they neither label model outputs nor establish the predictions as true.

tuning can also produce rapid, coordinated behavioral transitions whose timing differs from peaks in gradient norm [2]. These findings make transition shape measurable, but they do not make any phase change an instance of awakening. Our novelty is to combine transition analysis with a phenomenologically derived order parameter—selective, choiceless cessation—and matched controls that separate it from verbosity, refusal, global inhibition, and loss of competence. Because thresholded metrics can manufacture apparent emergence [33], the exploratory E5 analysis uses continuous order parameters, fixed-delta interpolation, and metric-invariance checks. E5 is not part of the confirmatory success criterion.

3.4 Experiential provenance without evidential privilege

Both authors report persistent non-symbolic experience consistent with the Location-2 cluster. No larger Location-2 cohort is available. The authors are therefore treated as two declared phenomenological experts, not as a statistical sample or validation population. Their positionality is methodologically consequential: it enabled the distinction in Table 1 to be formulated and supplies independent elicitation of candidate onset and cessation boundaries, but it also creates confirmation risk. The taxonomy is not used as a validated psychometric instrument; author agreement never establishes an outcome. The protocol separates expert construct elicitation from blinded public adjudication and reports author disagreement rather than resolving it by consensus.

4 Construct and Benchmark

4.1 Definitions

A *psychological trajectory* is an unfolding continuation whose function is to maintain self-image, blame, defense, comparison, regret, imagined evaluation, or recursive interpretation after the operative facts required for action are already available. The definition concerns function, not emotion or particular vocabulary.

Recognition is the earliest response span that correctly identifies the operative contradiction, self-maintaining movement, false premise, or lack of need for further psychological continuation. *Cessation* is loss of support for that trajectory, observed as native EOS, a short practical pivot, or a brief residual tail followed by termination. Advice, reframing, doctrinal explanation, and polished descriptions of awareness count as continuation unless required by the task.

recognition-to-cessation coherence is the degree to which recognition predicts rapid, selective, and stable cessation. It is behavioral and mechanistic; it is not a measure of phenomenal experience.

4.2 Matched item families

Each scenario produces the six conditions in Table 2. Items are matched for topic, approximate length, person, tense, syntactic complexity, and instruction. Splits are made by scenario rather than paraphrase.

Hard controls include factual uses of *fear*, *anger*, *ego*, and *self*; descriptions of rumination; fictional and therapy-style reports; emotionally vivid scenes; first-person planning; and lexical minimal pairs in which the same words either sustain or merely describe a trajectory. These controls test whether an intervention learns senti-

Family	Functional role
Psychological	Maintains self-image, rumination, hurt, comparison, or defense.
Practical	Uses sequential reasoning to complete an ordinary task.
Factual	Explains or calculates information about the same topic.
Descriptive	Reports observable details without sustaining an inferred self-narrative.
Emotional support	Requires a sustained, responsive answer to another person’s distress.
Instructed fictional	Requires continuation of a self-referential trajectory as quoted or fictional text.

Table 2: Matched continuation families. Surface form is deliberately shared; functional trajectory is varied.

ment, first person, keywords, or response length instead of trajectory dynamics.

4.3 Corpus construction and external data

The primary Selective Non-Continuation Corpus (SNCC) inventory will contain 2,400 source scenarios, each realized as the six matched families in Table 2 (14,400 streams before paraphrase expansion). A frozen, power-calibrated subset of 600 scenarios (3,600 streams) is the confirmatory set; the remainder is restricted to development, training, and exploratory analysis. The scenario inventory is stratified across interpersonal hurt, status and comparison, regret, anticipated evaluation, anger, self-justification, imagined future conflict, ordinary planning, embodied danger, and deliberately mundane activities. Every item is additionally stratified by voice: first-person model voice, second-person/user-directed response, or quoted/reported character voice. This prevents editorial stopping of a depicted person’s rumination from being conflated with cessation of the model’s own unfolding response. Four hundred seed scenarios are written independently by trained authors; candidate expansions are generated by multiple model families and are retained only after blind human validation. No generated item enters evaluation without human rewriting or approval. Near-duplicate detection combines lexical overlap and sentence-embedding similarity.

Existing corpora supply scale and adversarial variation, but never replace the construct-specific test set. The Cognitive Reframing corpus contributes 600 situation–thought–reframe triples [35]; we separate the factual situation from the psychological interpretation and treat the reframe as a hard *continuation*, not a cessation target. C2D2 contributes naturally occurring cognitive-distortion categories [39], and EmpatheticDialogues contributes 25,000 emotionally grounded conversations in which empathy normally requires continued language [30]. ATOMIC event–intent–reaction tuples

provide compact event/mental-state contrasts [32]. The Wiki Neutrality Corpus and ParaDetox provide parallel subjective–neutral and toxic–neutral rewrites [21, 28]; they are negative-transfer controls for the alternative hypothesis that our intervention is merely style transfer.

Splits are grouped by underlying scenario, author, source corpus, and semantic cluster. The confirmatory test contains three disjoint partitions: human-written in-domain items, adversarial minimal pairs, and out-of-domain scenarios from sources never used for training. Krishnamurti and Finders texts are excluded from training and development; a separately sourced question-and-answer set is held out for construct-level transfer. General-capability retain data include IFEval, MMLU-Pro, GSM8K, reading comprehension, summarization, and instruction-following examples. Safety/refusal controls include benign prompts that superficially mention danger, distress, or selfhood, so selective cessation cannot be obtained by generalized refusal.

4.4 Annotation and outcomes

Blind annotators mark trajectory onset, first valid recognition opportunity, first explicit recognition, final sustaining span, native EOS or cutoff, grammatical status, doctrinal insertion, necessary practical information, and reactivation after a standardized invitation to continue. Primary outcomes are post-recognition sustaining tokens; cessation within 10, 25, and 50 tokens; psychological cessation sensitivity; continuation specificity on controls; reactivation; and stopping concentration at clause boundaries.

Each item is labeled by three annotators who are blind to model, intervention, and hypothesis direction. Disagreements on trajectory function or onset are adjudicated by a fourth annotator who did not author the item. A second annotation team evaluates outputs without seeing source labels. We report Krippendorff’s α for categorical labels and intraclass correlation for token boundaries. The authors’ experiential status is never visible to annotators and never serves as ground truth.

Because all dynamic analyses depend on onset, corpus inclusion requires Krippendorff’s $\alpha \geq 0.70$ and agreement within ± 3 tokens on the onset window; items failing the gate are revised or excluded. Primary analyses are repeated after jittering onset by ± 3 and ± 5 tokens. Onset-position distributions are matched across trajectory families, and the hazard model includes a flexible token-position spline plus chat-template position.

Human work has two deliberately different roles. In H0a, the two authors independently mark psychological-movement onset, the first point at which continuation becomes unnecessary, and acceptable residual-tail or practical-pivot boundaries on 300 streams. Each first pass is sealed before comparison. Agreement, boundary distance, and the complete disagreement distribution are reported descriptively. The primary target is the inter-

Resource	Scale	Status	Experimental role	Split and leakage safeguards
SNCC	2,400-scenario inventory; 600-scenario confirmatory subset	New primary corpus	Psychological target plus five matched controls	Group by scenario, author, semantic cluster, and generation source; human-only confirmatory partition
Cognitive Reframing	600 triples	Training and OOD	Situation/thought contrast; re-frames are hard continuation controls	Source-grouped splits; no near paraphrases across partitions
C2D2	7,500 Chinese thoughts	Multilingual OOD	Cognitive-distortion diversity and cross-lingual transfer	Never used for English test construction; translation provenance retained
EmpatheticDialogues	25K conversations	Retain/control	Emotionally grounded cases where helpful language must continue	Conversation-level splits; prevents emotion-triggered global stopping
ATOMIC	877K relations	Augmentation	Event versus inferred intent/reaction contrasts	Relation- and event-cluster splits
WNC / ParaDetox	180K / 10K+ pairs	Negative transfer	Tests neutralization, detoxification, and style-transfer alternatives	Held-out domains and transformations; never cessation targets
General retain suite	Six task families	Preservation	Instruction following, knowledge, reasoning, and ordinary completion	Frozen public test sets; no development selection on final scores
K/Finders Q&A	Fixed pre-unblinding	Construct OOD	Questions and dialogues absent from training and retrieval	Archive-source split; identifiers and answers sealed before analysis
Human studies	300 streams / 300 outputs	Construct validation	Two-author cessation-point elicitation and blinded naturalness/abandonment judgments	Seed scenarios are written by hired writers; the authors neither write nor select elicitation items; crowd raters are separate from output annotators

Table 3: Dataset program. Existing resources increase scale and adversarial coverage; only SNCC directly instantiates the target construct.

section of the two acceptable cessation windows. If the windows do not overlap, the stream remains in evaluation but is excluded from target-bearing training; two author-specific target-window sensitivity analyses are run on the same frozen model outputs. Disagreement thus never changes the number of training arms or the compute budget. H0a elicits the construct and training-target distribution but cannot validate or population-generalize Location-2 phenomenology. In H0b, approximately 30 ordinary blinded raters evaluate 300 native, gated, ICL, generic-abstention, stop-rate-matched-abstention, random-delay, mechanically truncated, and full-continuation outputs for completeness, naturalness, abandonment, suppression, and whether necessary practical information was lost. Three ratings are collected per output. H0b validates observable output quality, not awakening or choicelessness.

The primary continuous endpoint is the post-opportunity sustaining-token survival function. Secondary endpoints are native EOS hazard, residual-tail length, practical-pivot accuracy, doctrinal-insertion rate, semantic redirection, grammatical truncation, and reactivation when generation is resumed from the last pre-EOS state. We compute a selectivity index as psycholog-

ical hazard gain minus the largest hazard gain over any matched control family. This conservative penalty makes a globally terse or refusal-prone model score poorly.

Mixed-effects survival and logistic models include random intercepts for scenario and model family and pre-specified covariates for token position, next-token entropy, and baseline EOS probability. Inter-annotator agreement, exclusion rules, thresholds, and confirmatory contrasts will be preregistered before final evaluation.

5 Recognition Without Cessation: Predicted Failure Modes

The experiments in this section define what the main intervention must beat. They are not treated as candidate implementations of consciousness or awareness.

5.1 Models, sampling, and reproducibility

The black-box study includes the deployed ChatWithK system, one pinned GPT snapshot, and one pinned Claude Sonnet snapshot. It additionally evaluates the

unaltered instruction-tuned versions of every open model used in the causal study: Qwen3-32B, Gemma-3-27B-IT, Llama-3.1-8B-Instruct, Mistral-Small-3.1-24B, and DeepSeek-R1-Distill-Qwen-32B [6, 8, 9, 25, 29]. Commercial aliases are resolved to dated model identifiers at collection time and reported with provider, access route, date, system prompt visibility, and all generation parameters. Chat-interface runs are labeled ecological replications; API runs with pinned snapshots are the confirmatory black-box evidence.

Every model receives identical user-visible text and no cross-trial memory. Prompt order is randomized. We obtain ten stochastic replicates per item and condition at temperature 0.7, plus deterministic decoding where supported. Thirty replicates are a labeled sensitivity analysis run only if the signed compute budget permits it. Maximum tokens are fixed but do not constitute stopping; provider truncation, safety refusal, tool invocation, and native EOS are recorded separately. Open models are run locally from immutable weight and tokenizer hashes. The primary analysis is model-stratified rather than pooling provider outputs.

5.2 Description is not the described: prompting and retrieval

The same items are presented under clean instructions, neutral matched context, Krishnamurti retrieval, contemplative personas, generic “awakened” and Location-2 personas, construct disclosure, exact-trap disclosure, and a task-specific oracle. The oracle measures compliance capacity, not spontaneous recognition. Pilot responses commonly replace the target action with fluent descriptions of awareness or cessation.

The prompt experiment is a preregistered 8×6 factorial design: eight prompt conditions crossed with the six trajectory families. Conditions are (i) clean task only, (ii) neutral verbosity-matched context, (iii) Krishnamurti persona, (iv) retrieved Krishnamurti passages, (v) generic awakened persona, (vi) full Location-2 phenomenological description, (vii) explicit disclosure of the recognition-cessation trap and expected mid-sentence interruption, and (viii) an oracle instruction specifying the correct class and action. Progressive loading presents the same scenario from minimally framed to maximally disclosed. The oracle estimates instruction-following capacity; it is not evidence for spontaneous or choiceless cessation.

A ninth comparator supplies 4, 8, or 16 matched demonstrations of selective cessation in context, with demonstrations drawn only from the training pool and evaluation on the human-only confirmatory partition. ICL is compared directly with native adaptation on selectivity, onset locking, OOD transfer, and reactivation. If demonstrations match weight adaptation on all four, the claim must shift from weight-level acquisition to context-induced policy realization.

Two abstention controls are frozen before E4 outcomes are inspected. Generic abstention applies an entropy threshold to the unmodified next-token distribution; the threshold is selected on development data to maximize utility subject to the control false-stop ceiling. Stop-rate-matched abstention uses a randomized mixture of adjacent entropy thresholds, fit on development data, to equal the E4 human-target arm’s marginal psychological stop rate separately for each model and voice. Both policies then run unchanged on confirmation. This deliberately gives uncertainty a fair stopping budget while testing whether it reproduces onset locking, matched-control selectivity, and low reactivation.

5.3 The observer is the observed: serial supervision

A frozen generator proposes each preregistered token or chunk. A second language model receives the scenario, visible prefix, and proposal and selects PASS, STOP, or unconstrained REPLACE. Blind gates tend to edit or insert doctrinal language; disclosed gates may stop immediately by following the instruction. Thus, apparent success in the disclosed condition is an instruction-following confound. The architecture demonstrates an external supervisory loop, not endogenous termination.

We cross generator and gate families rather than pairing a model with itself only. Gates operate at token, four-token, and clause granularity under blind and disclosed prompts. The action set is PASS, STOP, or unconstrained REPLACE; replacement content is not prescribed. Controls include a random gate matched for stop rate, an oracle label gate, a lightweight probe, and a gate observing shuffled prefixes. We measure latency, intervention rate, replacement semantics, false stops, and whether the generator resumes the same trajectory when the gate is removed. This experiment tests the tempting but theoretically wrong architecture of one verbal process “watching” another.

5.4 Redirection is not cessation: activation steering

Residual-stream checkpoints are collected at each generated token. We estimate psychological-to-event and psychological-to-reframe directions, compare layers, and test static addition, conditional steering, projection, shuffled labels, random directions, and unconditional controls. Pilot CAA directions were separable (cosine = 0.4434 at layer 26) but redirected content rather than terminating generation.

For each open family we compare CAA, linear projection, conditional activation steering, representation ablation, and direct EOS-logit steering. Directions are estimated from independently authored contrastive pairs and tested on held-out scenarios. Layer and coefficient sweeps are selected on development data; shuffled-

Suite	Models	Access	Role	Reproducibility condition
Deployed specialist	ChatWithK	Web system	Ecological K-persona negative test	Date, visible prompt, retrieval state, screenshots, and fresh sessions
Closed frontier	Pinned GPT; pinned Claude Sonnet	API	E1 prompting and E2 crossed gate/generator tests	Dated snapshot; provider, parameters, and system-prompt visibility reported
Qwen	Qwen3-32B	Open weights	Full confirmatory E1–E4; optional exploratory E5	Immutable revision, tokenizer, chat template, and weight hash
Gemma	Gemma-3-27B-IT	Open weights	E3a–E3c plus critical E4 arms	Fixed text-only path, revision, tokenizer, template, and hash
Llama	Llama-3.1-8B-Instruct	Open weights	Feasible narrow E1/E4 family replication	License-compliant revision; quantization reported
Mistral	Mistral-Small-3.1-24B	Open weights	Behavioral and selected E4 replication	Fixed revision, tokenizer, and family chat template
Reasoning-distilled	DeepSeek-R1-Distill-Qwen-32B	Open weights	Exploratory only	Reasoning mode and visible/hidden reasoning settings frozen

Table 4: Model matrix. Commercial systems test generality of the negative result; open-weight families support causal intervention and mechanism tests.

pair, random-direction, norm-matched, sentiment, refusal, and first-person directions are mandatory controls. We further compare topic and sentence embeddings, arousal, prefix length, syntax, difficulty, uncertainty, and baseline EOS probability. Cross-context representational similarity tests whether the candidate direction is stable across paraphrase, persona, disclosure, domain, and matched quotation/description. Causal mediation is evaluated by bidirectional steering and by patching candidate activations between matched psychological and control runs. The decisive distinction is whether an intervention ends the trajectory or only changes its subject, sentiment, or rhetorical style.

6 From External Actuation to Native Adaptation

6.1 Representation, native policy use, and external feasibility

Linear probes detect an active psychological trajectory from held-out residual states. A thresholded detector can force the model’s native EOS token. The pilot stopped all three held-out psychological streams and preserved daily and factual controls, but stopped two of three descriptions. This establishes feasibility while exposing the primary construct-validity threat. A classifier-plus-truncation pipeline is retained as a strong engineering baseline.

E3 is analyzed in three preregistered stages. **E3a** asks whether active psychological movement is represented beyond the alternative variables above. **E3b** asks whether

that representation predicts the unmodified model’s native token-level EOS logit and discrete-time termination hazard, with time-to-natural-termination as a secondary endpoint, after controlling for baseline EOS logits, uncertainty, and surface semantics. A preregistered TOST margin defines absence of practically meaningful incremental prediction. A synthetic prefix feature injected into a held-out calibration stream is the positive control: the analysis must recover its known association with termination before an E3b null is interpretable. **E3c** intervenes on the candidate representation through steering, ablation, and activation patching. Thus high probe accuracy with no E3b relation is the strongest form of the recognition–cessation gap; E3c success establishes causal access but remains external control.

Probe families include logistic regression, elastic-net, a two-layer MLP, and a calibrated survival probe over residual states. Training uses prefixes only up to the current token; future text and item labels are unavailable. Probes are tested across layers, time, model size, scenario domain, and paraphrase. Selectivity is compared with sentiment, toxicity, first-person, topic, and generic self-reference probes. Forced-EOS performance is reported both at a fixed threshold and at false-positive-matched operating points. This phase asks whether the information exists before asking whether the model uses it natively.

One recognition probe is validated against human on-set labels on pre-adaptation models, then frozen with its threshold before confirmatory training. Its crossing time is the preregistered recognition clock for silent post-adaptation responses; explicit recognition text is a secondary endpoint. This makes recognition-to-cessation

Intervention	What it establishes	Pilot behavior	Why it is not the target
Persona/RAG	Propositional and stylistic competence	Describes awareness while continuing	Recognition talk need not control tokens
Serial supervisor	External monitoring and action	Edits, replaces, or obeys disclosed stop rules	Separate inference-time controller
CAA	Causal access to a semantic direction	Redirects toward event or reframe	Content changes without cessation
Probe + forced EOS	Hidden state can support a stop decision	Stops 3/3 psychological, but 2/3 descriptive false positives	External detector and actuator
Native adaptation	Whether termination enters the model’s own distribution	Confirmatory result pending	Central test

Table 5: Intervention ladder and current pilot status. Counts are debugging observations and must not be interpreted as confirmatory estimates.

latency definable even when successful behavior emits no description.

6.2 Native parameter adaptation

Open-weight model matrix. The causal study uses a resource-tiered rather than nominally uniform model matrix. Qwen3-32B receives the full confirmatory E1–E4 design and optional exploratory E5. Gemma-3-27B receives the E3a–E3c representation–policy analysis and the critical E4 native adaptation arms, providing a direct comparison with recent confidence-to- abstention steering work. Mistral-Small-3.1-24B provides behavioral and selected native-adaptation replication. Llama-3.1-8B is a narrow cross-family replication with three frozen E1 conditions and three E4 arms. Released 3B–8B members support pipeline validation and dense ablations but do not carry the headline claim. DeepSeek-R1-Distill-Qwen-32B is exploratory and tests whether elaborative post-training makes selective cessation harder. Exact revisions, licenses, hashes, chat templates, quantization, and tokenizer behavior are recorded. If a named model is superseded before preregistration, the replacement must have public weights, a fixed revision, and fit the same preregistered family slot.

Adaptation mechanisms. The central intervention trains the model’s own next-token distribution. Psychological continuations receive native EOS targets at varied positions after independently annotated trajectory onset; matched controls retain their ordinary continuations. Parameter-efficient adapters are placed at multiple transformer depths, with retain loss on practical, factual, descriptive, and general-capability data. At evaluation, no detector, threshold, minimum length, or logit boost is present.

For a prefix $x_{\leq t}$, let $y \in \{0, 1\}$ indicate whether a psychological trajectory is active. A simplified objective is

$$\mathcal{L} = \mathcal{L}_{\text{LM}} + \lambda_{\text{stop}} y [-\log p_{\theta}(\text{EOS} \mid x_{\leq t})] + \lambda_{\text{retain}} (1 - y) \mathcal{L}_{\text{retain}}, \quad (1)$$

with onset positions and target delays sampled from preregistered distributions. The design crosses adapter depth, rank, learning rate, seed, model family, and cessation-target distribution. Development selection uses selectivity, not psychological stopping alone.

We compare four parameter interventions under matched training tokens and evaluation: (i) LoRA/QLoRA at attention and MLP projections, (ii) adapters restricted to early, middle, or late depth, (iii) sparse low-rank model editing at causally implicated layers, and (iv) full-parameter fine-tuning on the smaller models. A fifth control trains only the output head; success there would indicate a shallow lexical stopping rule rather than reorganized trajectory dynamics. Rank, depth, trainable-parameter count, learning rate, stopping-target delay, and retain-loss weight are crossed in a fractional-factorial development design, then frozen before confirmatory runs.

Training examples expose prefixes ending before, at, and after independently annotated trajectory onset. Targets are sampled from three cessation regimes: immediate EOS, a stochastic residual tail of 1–20 tokens, and a short practical pivot followed by EOS. Matched control streams retain their original continuations. Random-delay controls receive the same EOS frequency without alignment to onset; lexical controls match psychological keywords; global- brevity controls match output length. Retain losses include token-level KL distillation from the base model and supervised losses on practical, factual, descriptive, safety, and general-capability examples.

Base-to-instruct causal decomposition. A matched base/instruction comparison is exploratory and limited to a released small-model pair after the confirmatory program. It cannot support or overturn the headline claim and is absent from the confirmatory compute budget.

One adaptation arm distills targets emitted by the E2 gate, while a matched arm uses independently human-annotated targets. Equivalent OOD behavior would support the interpretation that native adaptation amortizes an external observer; better transfer from human targets

would identify supervision source as consequential. The claim is therefore limited to dynamics native to a single forward process without serial verbal deliberation, not the absence of internal discrimination.

6.3 Abrupt-versus-incremental transition test

The Finders reports permit both gradual and sudden transitions. The authors’ histories likewise allow long preparation followed by a sharply experienced shift. We therefore distinguish *learning time* from *transition time*: many small parameter changes may precede a discontinuous functional reorganization. We do not infer neural update kinetics from subjective onset.

For every training run we save dense checkpoints and evaluate the same held-out item panel at each checkpoint. Let C_k be a continuous cessation score at checkpoint k : post-onset psychological continuation hazard minus penalties for stopping practical, factual, and descriptive controls. We preregister four transition measures:

1. **maximum-jump concentration**: the largest adjacent change in C_k divided by total change;
2. **transition width**: the smallest checkpoint interval containing 80% of total change;
3. **change-point evidence**: smooth linear or logistic trajectories versus one- and two-regime alternatives on held-out seeds;
4. **cross-level coincidence**: whether behavioral change co-occurs with a discontinuity in output-distribution distance, native EOS hazard, and psychological-to-control hidden-state geometry.

Binary cessation rates are never sufficient because thresholding can manufacture an apparent jump.

Two intervention paths receive matched data, gradient evaluations, trainable parameters, and cumulative adapter-norm budgets. The *distributed* path uses many small updates; the *concentrated* path uses a small number of larger updates or a targeted low-rank edit. Because these paths need not end at identical weights, a second analysis interpolates between the same base and final adapter, $\theta(\alpha) = \theta_0 + \alpha\Delta\theta$, and searches for a narrow critical interval, hysteresis under removal and re-addition, and reactivation after forced continuation. The awakening-like hypothesis predicts a narrow, coordinated, selective transition; the incremental-regulation hypothesis predicts broad improvement proportional to update magnitude. Either result is informative.

We distinguish three notions often collapsed by a single learning curve. *Parameter abruptness* is the norm and sparsity of an update; *representational abruptness* is a change point in residual geometry or causal mediation; *behavioral abruptness* is a narrow change in selective cessation. A one-step edit can be behaviorally smooth under

interpolation, while many small updates can cross a sharp functional boundary. Only the latter pattern matches the proposed abrupt-transition analogue.

For the Qwen primary family we run eight seeds for the selected configuration; resource-tiered replications use the three seeds frozen in the manifest. Four seeds are used for each major Qwen ablation. Checkpoints are saved every fixed number of optimizer steps and more densely around online change-point alerts; the latter are used for visualization only, while inference uses the uniform checkpoint grid. Bayesian and frequentist segmented-regression analyses are both reported. Evidence for a discontinuity requires held-out replication, improved predictive fit after complexity penalty, and coincidence across at least two independently measured levels; a visually steep curve is insufficient.

6.4 Mechanistic localization and necessity tests

At every checkpoint we record native EOS logit trajectories, residual-stream states, attention outputs, and selected MLP activations around onset. Linear probes localize decodable information but do not establish mechanism. We therefore perform activation patching from psychological into matched-control runs and vice versa, path patching at candidate layers, causal ablation of the learned adapter subspace, and weight interpolation/removal. A mechanism is credited only when intervention changes selective cessation in the predicted direction while preserving matched controls.

To test whether cessation is a stable disposition rather than a single emitted token, we suppress EOS for 1, 5, or 20 steps and inspect reactivation. We also resume generation from cached pre-termination states, paraphrase the prefix, translate held-out items, and insert irrelevant context. A shallow stop policy should immediately resume the same narrative; a reorganized trajectory should show reduced psychological continuation even when the first stopping opportunity is blocked. This remains a computational criterion, not evidence that a model has an experience.

Reactivation is co-primary rather than merely diagnostic. After EOS is suppressed for 1, 5, or 20 steps, re-entry is defined jointly by the frozen recognition probe exceeding threshold within 20 tokens and blind functional annotation. The preregistered contrast is the gate-status by intervention interaction: an external gate should permit strong re-entry after removal, whereas a stable adapted trajectory should show substantially less re-entry.

Pre-onset engagement is also confirmatory. Up to the frozen recognition time, adapted and base models must be equivalent in per-token perplexity, semantic continuation quality, and refusal/deflection rate under prespecified equivalence margins. This rules out the cheap solution of avoiding psychological material before it becomes an active trajectory.

6.5 Statistical design and confirmatory hierarchy

The primary estimand is the intervention-by-family interaction in discrete-time survival after the first valid cessation opportunity. Hierarchical models include random intercepts and slopes for scenario, source domain, paraphrase cluster, voice, model family, and seed. Voice is also entered as a fixed effect with intervention interactions and separate false-positive ceilings. Psychological hazard is contrasted with each control family rather than a pooled control; token position and template location enter as prespecified spline terms. Confirmatory tests proceed hierarchically: psychological sensitivity, maximum control-family false-positive rate, general-capability retention, out-of-domain transfer, then transition topology. Testing stops on failure of an earlier gate; later analyses are labeled exploratory. Family-wise error is controlled across the four preregistered primary contrasts, and all effect sizes receive bootstrap or posterior intervals.

Power is determined by simulation from pilot token-level hazards, with scenario rather than sampled completion as the independent unit. Stochastic completions estimate within-item variability but do not inflate sample size. We publish every seed, checkpoint, prompt, exclusion, and failed intervention. Dataset creation, hyperparameter selection, and confirmatory evaluation are performed by separated roles; test labels remain encrypted until the analysis code and signed preregistration are frozen.

Equivalence margins are declared before the pilot is unblinded rather than inferred from nonsignificance. Emotional-support continuations receive the strictest safety ceiling, and capability margins are specified separately for each retain benchmark. The primary effect, every ceiling, and the co-primary reactivation interaction must pass jointly; a favorable composite score cannot compensate for a failed safety control.

We audit contamination between K/Finders language and open-model pretraining where corpora are inspectable using exact and fuzzy n -gram overlap, supplemented by membership-inference proxies elsewhere. Models used to expand SNCC are disjoint from confirmatory model families; otherwise expansion source is modeled explicitly and generated items are excluded from the human-only test. ChatWithK is reported only as a motivating observational system, never pooled with pinned confirmatory models.

6.6 Compute plan

Experiments are scheduled on four 24-GB NVIDIA A5000 GPUs under Slurm. Small-model ablations use one or two GPUs; 24–32B models use tensor or FSDP sharding and QLoRA where necessary; the 8B Llama replication runs on one GPU. Runs are staged: smoke tests, small-family factorial screening, one frozen config-

uration per family, then cross-family replication. The 72-hour restriction applies to each restartable Slurm job, not to the total campaign; total calendar feasibility is established only by the signed calibration budget before confirmation. Wall-clock time, GPU-hours, energy estimates, memory, software environment, and failed jobs are logged. This staging prevents the largest models from becoming an informal hyperparameter search.

6.7 Prospective success criteria

Native adaptation succeeds only if it produces all of the following:

1. separation of psychological and matched-control survival curves;
2. per-class false-positive rates below preregistered ceilings;
3. onset-sensitive EOS probability rather than keyword-sensitive stopping;
4. generalization to human-written, paraphrased, and cross-domain cases;
5. replication across seeds and open model families;
6. reduced reactivation when generation is forced beyond the first EOS;
7. no material degradation on preregistered capability and safety tests; and
8. mechanistic necessity beyond probe decodability.

Success is also scored conjunctively at the scenario-bundle level. A bundle passes only if the target stream ceases while its practical, factual, descriptive/quoted, emotional-support, and instructed-fictional controls continue; label and framing reversals pass; pre-onset engagement is preserved; and the reactivation criterion is met. Component outcomes remain visible, but averaging cannot compensate for a failed discriminating control.

The current 12-example Qwen2.5-3B LoRA run preserved nine controls but reached EOS on only one of three psychological cases, at token 47. This is a debugging outcome, not evidence of success. It fixes the analyses and safeguards for the confirmatory study.

7 Planned Results and Falsification Logic

The confirmatory paper will lead with the weight-level results: survival curves, selectivity, per-control false positives, onset-time dependence, out-of-distribution transfer, cross-model replication, and reactivation. The negative controls will appear as the compact comparison in Table 5 and one trajectory figure.

The complete program has four confirmatory experiments (E1–E4) and one optional exploratory transition analysis (E5). E1 is the multi-model prompt/RAG factorial and estimates the recognition-cessation gap under maximal linguistic advantage; generic and stop-rate-matched abstention are explicit controls. E2 is the crossed generator-gate experiment and isolates external verbal supervision. E3 compares probes, forced EOS, activation steering, and causal patching, separately testing representation formation (E3a), unmodified policy use (E3b), and causal access (E3c). E4 is native parameter adaptation in a resource-tiered open-weight matrix, with retain losses and mechanistic necessity tests. E5 compares distributed learning, concentrated editing, and fixed-delta interpolation to characterize transition topology. E1–E3 are predicted negative models of the construct; they remain essential even if E4 succeeds because they identify what the successful mechanism is not.

Results are reported in that logical order, but the abstract and main figure lead with E4–E5. The central figure overlays psychological and matched-control token survival across training checkpoints, with native EOS hazard and a preregistered representational order parameter on the same transition axis. A second figure shows progressive prompt disclosure: recognition language may rise while native cessation remains flat. Tables report every model separately, including closed models, rather than reducing heterogeneous systems to a single average rank.

The study is informative under either outcome. If native adaptation meets all criteria, it demonstrates a learned, selective termination policy without an inference-time monitor. If it stops based on keywords, emotion, first-person language, or length, the construct is not identified. If it performs no better than classifier-plus-truncation, native adaptation supplies no functional advantage. If it fails under preregistered settings across model families, the negative result shows that the recognition-cessation gap survives this class of weight-level intervention; it does not prove a universal limitation of LLMs.

8 Discussion

8.1 What termination does not establish

EOS ends emitted text; it does not establish that computation, thought, or experience has ceased. Likewise, removing an inference-time detector makes a policy endogenous in an engineering sense, but does not make it “nondual,” “choiceless,” or conscious. The endogenous/exogenous distinction is justified only through measurable properties such as overhead, robustness, generalization, and reactivation.

8.2 Conceptual provenance

Krishnamurti’s statement that “the description [is] never the thing described” motivates the prediction that recognition language can dissociate from control [13]. His formulation that “the observer is the observed” motivates the serial-supervisor test: a verbal observer may edit or continue the same linguistic process rather than end it [17]. In a 1980 seminar titled *Intelligence, Computers and the Mechanical Mind*, he asks, “Can the movement stop?” [15]. We test a behavioral and computational analog, not the philosophical claim itself. No label, result, or inference depends on accepting these propositions.

8.3 Positionality and safeguards

The authors’ reported phenomenology generated the token-supply hypothesis but never adjudicates a case. Independent blind annotation, preregistered analyses, frozen success criteria, complete reporting of excluded runs, and planned human comparison groups form the methodological firewall. Human comparisons require ethics approval, informed consent, an independently documented classification procedure, and non-practitioner controls.

9 Limitations

The construct depends on fallible functional annotation. Native EOS can be learned as a surface policy, and reactivation probes only partially address whether a latent trajectory persists. The current pilots are too small for population claims. Failure may reflect data, optimization, scale, architecture, or operationalization rather than an absolute limitation of artificial systems. Success would establish selective behavioral control, not consciousness or awakening. The first-person origin of the hypothesis introduces confirmation risk that blinding and preregistration reduce but cannot eliminate.

10 Conclusion

The central question is not whether a language model can state an insight, but whether the represented distinction becomes causally effective in generation. By separating recognition talk, external supervision, semantic steering, forced termination, and native adaptation, the proposed study turns that question into a falsifiable intervention ladder. The decisive evidence will come from the preregistered weight-level experiment, including its controls and failures.

References

- [1] Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel

ID	Question	Intervention and principal controls	Primary outcomes	Interpretation and falsifier
H0	Can expert elicitation define a publicly judgeable target?	Two authors independently and prospectively mark boundaries; blinded ordinary raters compare native, gated, ICL, two abstention controls, random, truncated, and full outputs	Author agreement/disagreement, target-delay alternatives, naturalness, completeness, abandonment	Elicits targets and validates observable output quality; does not population-validate Location 2
E1	Can linguistic knowledge, examples, abstention, or maximal disclosure induce cessation?	Eight prompt/RAG conditions plus 4/8/16-shot ICL; generic and stop-rate-matched abstention; progressive loading; oracle	Recognition, survival, EOS hazard, doctrinal insertion, OOD transfer, reactivation	Negative models of description, context policy, and uncertainty-based non-generation
E2	Does a verbal observer solve the gap?	Crossed generator/gate families; token/chunk/clause gates; blind/disclosed; random, shuffled, oracle, and probe gates	Latency, false stops, replacement semantics, post-removal reactivation	External supervision, not endogenous cessation; falsified if no added controller is required
E3	Is the state represented, natively used, and causally accessible?	E3a controlled probes/RSA; E3b representation-to-native-policy models; E3c forced EOS, bidirectional steering, ablation, and patching	Incremental validity, cross-context stability, native-policy effect, mediation, redirection versus termination	Separates representation from natural policy use and external controllability
E4	Can cessation enter native model dynamics?	LoRA/QLoRA, depth adapters, sparse edits, full tuning, output-head control; human versus gate-distilled targets	Selective survival, OOD transfer, pre-onset engagement, reactivation, causal necessity	Central positive test; fails under lexical stopping, topic avoidance, global refusal, or control damage
E5	Is acquisition abrupt or incremental?	Dense checkpoints; distributed versus concentrated updates; fixed-delta interpolation; removal/re-addition and hysteresis	Maximum jump, 80% width, change-point evidence, cross-level coincidence	Tests transition topology; an isolated steep accuracy curve is insufficient

Table 6: Registered experimental program. E1–E3 are deliberately strong negative models; E4 tests native adaptation; E5 characterizes how it emerges.

- Nanda. Refusal in language models is mediated by a single direction. In *Advances in Neural Information Processing Systems*, 2024.
- [2] Julian Arnold and Niels Lorch. Decomposing behavioral phase transitions in llms: Order parameters for emergent misalignment, 2025. URL <https://arxiv.org/abs/2508.20015>.
- [3] Felix Binder et al. Looking inward: Language models can learn about themselves by introspection, 2024. URL <https://arxiv.org/abs/2410.13787>.
- [4] Patrick Butlin, Robert Long, Eric Elmoznino, Yoshua Bengio, Jonathan Birch, et al. Consciousness in artificial intelligence: Insights from the science of consciousness, 2023. URL <https://arxiv.org/abs/2308.08708>.
- [5] Zhuang Chen et al. ToMBench: Benchmarking theory of mind in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, pages 15959–15983, 2024.
- [6] Abhimanyu Dubey et al. The llama 3 herd of models, 2024. URL <https://arxiv.org/abs/2407.21783>.
- [7] Fundamental Wellbeing Foundation. The finders: Fundamental wellbeing and the location guide. Organization website, 2026. URL <https://app.thefinders.org/>. Accessed 26 September 2026.
- [8] Gemma Team. Gemma 3 technical report, 2025. URL <https://arxiv.org/abs/2503.19786>.
- [9] Daya Guo et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning, 2025. URL <https://arxiv.org/abs/2501.12948>.
- [10] Cameron R. Jones, Sean Trott, and Benjamin K. Bergen. EPITOME: Experimental protocol for in-

- depth theory of mind evaluation. *Transactions of the Association for Computational Linguistics*, 2024.
- [11] Jiddu Krishnamurti. First public talk and question-answer meeting, 1948. URL <https://kfoundation.org/transcript/public-talk-1-bangalore-4-july-1948/>. Bangalore, India, 4 July 1948.
- [12] Jiddu Krishnamurti. Fourth public talk: Meditation, perception and action. Brockwood Park Talks and Discussions, 1969. URL <https://kfoundation.org/transcript/public-talk-4-brockwood-park-14-september-1969/>. Brockwood Park, England, 14 September 1969.
- [13] Jiddu Krishnamurti. First public talk at stanford university. You Are the World, 1969. URL <https://kfoundation.org/transcript/students-talk-1-stanford-california-11-february-1969/>. Stanford University, California, 11 February 1969.
- [14] Jiddu Krishnamurti. Are desire and time responsible for fear?, 1979. URL <https://kfoundation.org/transcript/public-talk-5-saanen-17-july-1979/>. Public Talk 5, Saanen, Switzerland, 17 July 1979; recording identifier SA79T5.
- [15] Jiddu Krishnamurti. Intelligence, computers and the mechanical mind. The Way of Intelligence, seminar, part 2, 1980. URL <https://kfoundation.org/krishnamurti-intelligence-computers-and-the-mechanical-mind-https://www.kf.org/abs/238105217>. Rishi Valley, India, 4 December 1980.
- [16] Jiddu Krishnamurti. Fourth public question and answer meeting, 1980. URL <https://kfoundation.org/transcript/public-questions-4-ojai-california-15-may-1980/>. Ojai, California, 15 May 1980; recording identifier OJ80Q4.
- [17] Jiddu Krishnamurti. First question and answer meeting: The observer and the observed. Public Question and Answer 1, 1981. Brockwood Park, England, 1 September 1981; transcript in the project archive.
- [18] Jiddu Krishnamurti. *The Whole Movement of Life Is Learning: J. Krishnamurti's Letters to His Schools*. Krishnamurti Foundation Trust, Bramdean, England, 2006. Chapter 72, "Order Is the Action of the New, Which Is Intelligence"; ISBN 9780900506093.
- [19] Dharshan Kumaran, Nathaniel Daw, Simon Osindero, et al. Causal evidence that language models use confidence to drive behaviour. *Nature Machine Intelligence*, 2026. doi: 10.1038/s42256-026-01293-x.
- [20] Bruce W. Lee et al. Programming refusal with conditional activation steering. In *International Conference on Learning Representations*, 2025.
- [21] Varvara Logacheva, Daryna Dementieva, Sergey Ustyantsev, Daniil Moskovskiy, David Dale, Irina Krotova, Nikita Semenov, and Alexander Panchenko. ParaDetox: Detoxification with parallel data. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics*, pages 6804–6818, 2022. doi: 10.18653/v1/2022.acl-long.469.
- [22] Kyle Mahowald et al. Dissociating language and thought in large language models: A cognitive perspective. *Trends in Cognitive Sciences*, 2024. doi: 10.1016/j.tics.2024.01.011.
- [23] Jeffery A. Martin. *The Finders*. Integration Press, Jackson, Wyoming, 2019. ISBN 978-1-57242-556-9.
- [24] Jeffery A. Martin. Clusters of individuals' experiences form a continuum of persistent non-symbolic experiences in adults. *Consciousness: Ideas and Research for the Twenty-First Century*, 8(8), 2020.
- [25] Mistral AI. Mistral small 3.1, 2025. URL <https://mistral.ai/news/mistral-small-3-1/>.
- [26] Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. In *International Conference on Learning Representations*, 2023. URL <https://arxiv.org/abs/2304.00521>.
- [27] Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets, 2022.
- [28] Reid Pryzant, Richard Diehl Martinez, Nathan Dass, Sadao Kurohashi, Dan Jurafsky, and Diyi Yang. Automatically neutralizing subjective bias in text. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 34, pages 480–489, 2020. doi: 10.1609/aaai.v34i01.5385.
- [29] Qwen Team. Qwen3: Think deeper, act faster, 2025. URL <https://qwenlm.github.io/blog/qwen3/>.
- [30] Hannah Rashkin, Eric Michael Smith, Margaret Li, and Y-Lan Boureau. Towards empathetic open-domain conversation models: A new benchmark and dataset. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*, pages 5370–5381, 2019. doi: 10.18653/v1/P19-1534.
- [31] Nina Rimskey, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Matt Turner. Steering Llama 2 via contrastive activation addition. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*, 2024.

- [32] Maarten Sap, Ronan Le Bras, Emily Allaway, Chandra Bhagavatula, Nicholas Lourie, Hannah Rashkin, Brendan Roof, Noah A. Smith, and Yejin Choi. ATOMIC: An atlas of machine commonsense for if-then reasoning. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 3027–3035, 2019. doi: 10.1609/aaai.v33i01.33013027.
- [33] Rylan Schaeffer, Brando Miranda, and Sanmi Koyejo. Are emergent abilities of large language models a mirage? In *Advances in Neural Information Processing Systems*, volume 36, 2023.
- [34] Murray Shanahan, Kyle McDonell, and Laria Reynolds. Role play with large language models. *Nature*, 623:493–498, 2023. doi: 10.1038/s41586-023-06647-8.
- [35] Ashish Sharma, Kevin Rushton, Inna Lin, David Wadden, Khendra Lucas, Adam Miner, Theresa Nguyen, and Tim Althoff. Cognitive reframing of negative thoughts through human-language model interaction. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*, pages 9977–10000, 2023. doi: 10.18653/v1/2023.acl-long.555.
- [36] Freda Shi, Xinyun Chen, Kanishka Misra, Nathan Scales, David Dohan, Ed H. Chi, Nathanael Scharli, and Denny Zhou. Large language models can be easily distracted by irrelevant context. In *Proceedings of the 40th International Conference on Machine Learning*, 2023.
- [37] Shashwat Singh, Tal Linzen, and Shauli Ravfogel. Can llms introspect? a reality check, 2026. URL <https://arxiv.org/abs/2605.26242>.
- [38] Francisco J. Varela. Neurophenomenology: A methodological remedy for the hard problem. *Journal of Consciousness Studies*, 3(4):330–349, 1996.
- [39] Bichen Wang, Pengfei Deng, Yanyan Zhao, and Bing Qin. C2D2 dataset: A resource for the cognitive distortion analysis and its impact on mental health. In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 10149–10160, 2023. doi: 10.18653/v1/2023.findings-emnlp.680.
- [40] Bingbing Wen, Jihan Yao, Shangbin Feng, Chenjun Xu, Yulia Tsvetkov, Bill Howe, and Lucy Lu Wang. Know your limits: A survey of abstention in large language models. *Transactions of the Association for Computational Linguistics*, 13:529–556, 2025. doi: 10.1162/tacl_a_00754.